

PS4

January 15, 2024

0.1 Data organization

0.2 Data loading

```
[1]: pip install shap
```

```
Requirement already satisfied: shap in c:\programdata\anaconda3\lib\site-  
packages (0.41.0)  
Requirement already satisfied: slicer==0.0.7 in  
c:\programdata\anaconda3\lib\site-packages (from shap) (0.0.7)  
Requirement already satisfied: packaging>20.9 in  
c:\programdata\anaconda3\lib\site-packages (from shap) (21.0)  
Requirement already satisfied: scipy in c:\programdata\anaconda3\lib\site-  
packages (from shap) (1.7.1)  
Requirement already satisfied: numpy in c:\programdata\anaconda3\lib\site-  
packages (from shap) (1.20.3)  
Requirement already satisfied: numba in c:\programdata\anaconda3\lib\site-  
packages (from shap) (0.54.1)  
Requirement already satisfied: pandas in c:\programdata\anaconda3\lib\site-  
packages (from shap) (1.3.4)  
Requirement already satisfied: cloudpickle in c:\programdata\anaconda3\lib\site-  
packages (from shap) (2.0.0)  
Requirement already satisfied: scikit-learn in  
c:\programdata\anaconda3\lib\site-packages (from shap) (0.24.2)  
Requirement already satisfied: tqdm>4.25.0 in c:\programdata\anaconda3\lib\site-  
packages (from shap) (4.62.3)  
Requirement already satisfied: pyparsing>=2.0.2 in  
c:\programdata\anaconda3\lib\site-packages (from packaging>20.9->shap) (3.0.4)  
Requirement already satisfied: colorama in c:\programdata\anaconda3\lib\site-  
packages (from tqdm>4.25.0->shap) (0.4.4)  
Requirement already satisfied: llvmlite<0.38,>=0.37.0rc1 in  
c:\programdata\anaconda3\lib\site-packages (from numba->shap) (0.37.0)  
Requirement already satisfied: setuptools in c:\programdata\anaconda3\lib\site-  
packages (from numba->shap) (58.0.4)  
Requirement already satisfied: python-dateutil>=2.7.3 in  
c:\programdata\anaconda3\lib\site-packages (from pandas->shap) (2.8.2)  
Requirement already satisfied: pytz>=2017.3 in  
c:\programdata\anaconda3\lib\site-packages (from pandas->shap) (2021.3)  
Requirement already satisfied: six>=1.5 in c:\programdata\anaconda3\lib\site-
```

```

packages (from python-dateutil>=2.7.3->pandas->shap) (1.16.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in
c:\programdata\anaconda3\lib\site-packages (from scikit-learn->shap) (2.2.0)
Requirement already satisfied: joblib>=0.11 in
c:\programdata\anaconda3\lib\site-packages (from scikit-learn->shap) (1.1.0)
Note: you may need to restart the kernel to use updated packages.

```

```

WARNING: Ignoring invalid distribution -cikit-learn
(c:\programdata\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -cikit-learn
(c:\programdata\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -cikit-learn
(c:\programdata\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -cikit-learn
(c:\programdata\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -cikit-learn
(c:\programdata\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -cikit-learn
(c:\programdata\anaconda3\lib\site-packages)

```

```
[1]: from shap.datasets import adult
```

```

X, y = adult()
print(X.shape)
print(y.shape)
print(type(X))
print(type(y))

```

```

(32561, 12)
(32561,)
<class 'pandas.core.frame.DataFrame'>
<class 'numpy.ndarray'>

```

```
[2]: print(X)
print(y)
```

	Age	Workclass	Education-Num	Marital Status	Occupation	\
0	39.0	7	13.0	4	1	
1	50.0	6	13.0	2	4	
2	38.0	4	9.0	0	6	
3	53.0	4	7.0	2	6	
4	28.0	4	13.0	2	10	
...	
32556	27.0	4	12.0	2	13	
32557	40.0	4	9.0	2	7	
32558	58.0	4	9.0	6	1	
32559	22.0	4	9.0	4	1	
32560	52.0	5	9.0	2	4	

	Relationship	Race	Sex	Capital Gain	Capital Loss	Hours per week \
0	0	4	1	2174.0	0.0	40.0
1	4	4	1	0.0	0.0	13.0
2	0	4	1	0.0	0.0	40.0
3	4	2	1	0.0	0.0	40.0
4	5	2	0	0.0	0.0	40.0
...
32556	5	4	0	0.0	0.0	38.0
32557	4	4	1	0.0	0.0	40.0
32558	1	4	0	0.0	0.0	40.0
32559	3	4	1	0.0	0.0	20.0
32560	5	4	0	15024.0	0.0	40.0

	Country
0	39
1	39
2	39
3	39
4	5
...	...
32556	39
32557	39
32558	39
32559	39
32560	39

```
[32561 rows x 12 columns]
[False False False ... False False  True]
```

```
[3]: print(X['Age'])
      print(X['Age'].shape)
      print(X['Age'].values)
      print(type(X['Age'].values))
      print(X['Age'].values.reshape(-1,1))
      print(X['Age'].values.reshape(-1,1).shape)
```

0	39.0
1	50.0
2	38.0
3	53.0
4	28.0
...	
32556	27.0
32557	40.0
32558	58.0
32559	22.0
32560	52.0

```
Name: Age, Length: 32561, dtype: float32
```

```
(32561,)
[39. 50. 38. ... 58. 22. 52.]
<class 'numpy.ndarray'>
[[39.]
 [50.]
 [38.]
 ...
 [58.]
 [22.]
 [52.]]
(32561, 1)
```

```
[4]: numerical_columns = ['Age', 'Education-Num', 'Capital Gain', 'Capital Loss', 'Hours_
    ↪per week']
    categorical_columns = ['Workclass', 'Marital_
    ↪Status', 'Occupation', 'Relationship', 'Race', 'Sex', 'Country']
```

0.3 Conversion of categorical data

```
[5]: import pandas as pd # for one-hot encoding
    from sklearn.preprocessing import StandardScaler # for normalization
```

```
[6]: # Normalization of numerical data
    for column in numerical_columns:
        scaler = StandardScaler()
        X[column] = scaler.fit_transform(X[column].values.reshape(-1,1))
        #X[column] = scaler.fit_transform(X[column])

    print(X)
    print(type(X))
```

	Age	Workclass	Education-Num	Marital Status	Occupation	\
0	0.030671	7	1.134739	4	1	
1	0.837109	6	1.134739	2	4	
2	-0.042642	4	-0.420060	0	6	
3	1.057047	4	-1.197459	2	6	
4	-0.775768	4	1.134739	2	10	
...	
32556	-0.849080	4	0.746039	2	13	
32557	0.103983	4	-0.420060	2	7	
32558	1.423610	4	-0.420060	6	1	
32559	-1.215643	4	-0.420060	4	1	
32560	0.983734	5	-0.420060	2	4	

	Relationship	Race	Sex	Capital Gain	Capital Loss	Hours per week	\
0	0	4	1	0.148453	-0.21666	-0.035429	
1	4	4	1	-0.145920	-0.21666	-2.222153	

2	0	4	1	-0.145920	-0.21666	-0.035429
3	4	2	1	-0.145920	-0.21666	-0.035429
4	5	2	0	-0.145920	-0.21666	-0.035429
...
32556	5	4	0	-0.145920	-0.21666	-0.197409
32557	4	4	1	-0.145920	-0.21666	-0.035429
32558	1	4	0	-0.145920	-0.21666	-0.035429
32559	3	4	1	-0.145920	-0.21666	-1.655225
32560	5	4	0	1.888424	-0.21666	-0.035429

	Country
0	39
1	39
2	39
3	39
4	5
...	...
32556	39
32557	39
32558	39
32559	39
32560	39

```
[32561 rows x 12 columns]
<class 'pandas.core.frame.DataFrame'>
```

```
[7]: # Data type change of categorical data
# categorical data

for column in categorical_columns:
    X[column] = X[column].astype('category')

print(X)
print(X['Country'].values)
```

	Age	Workclass	Education-Num	Marital Status	Occupation	\
0	0.030671	7	1.134739		4	1
1	0.837109	6	1.134739		2	4
2	-0.042642	4	-0.420060		0	6
3	1.057047	4	-1.197459		2	6
4	-0.775768	4	1.134739		2	10
...
32556	-0.849080	4	0.746039		2	13
32557	0.103983	4	-0.420060		2	7
32558	1.423610	4	-0.420060		6	1
32559	-1.215643	4	-0.420060		4	1
32560	0.983734	5	-0.420060		2	4

	Relationship	Race	Sex	Capital Gain	Capital Loss	Hours per week	\
0		0	4	1	0.148453	-0.21666	-0.035429
1		4	4	1	-0.145920	-0.21666	-2.222153
2		0	4	1	-0.145920	-0.21666	-0.035429
3		4	2	1	-0.145920	-0.21666	-0.035429
4		5	2	0	-0.145920	-0.21666	-0.035429
...
32556		5	4	0	-0.145920	-0.21666	-0.197409
32557		4	4	1	-0.145920	-0.21666	-0.035429
32558		1	4	0	-0.145920	-0.21666	-0.035429
32559		3	4	1	-0.145920	-0.21666	-1.655225
32560		5	4	0	1.888424	-0.21666	-0.035429

	Country
0	39
1	39
2	39
3	39
4	5
...	...
32556	39
32557	39
32558	39
32559	39
32560	39

[32561 rows x 12 columns]

[39, 39, 39, 39, 5, ..., 39, 39, 39, 39, 39]

Length: 32561

Categories (42, int64): [0, 1, 2, 3, ..., 38, 39, 40, 41]

```
[8]: X = pd.get_dummies(X)
      print(X)
```

	Age	Education-Num	Capital Gain	Capital Loss	Hours per week	\
0	0.030671	1.134739	0.148453	-0.21666	-0.035429	
1	0.837109	1.134739	-0.145920	-0.21666	-2.222153	
2	-0.042642	-0.420060	-0.145920	-0.21666	-0.035429	
3	1.057047	-1.197459	-0.145920	-0.21666	-0.035429	
4	-0.775768	1.134739	-0.145920	-0.21666	-0.035429	
...
32556	-0.849080	0.746039	-0.145920	-0.21666	-0.197409	
32557	0.103983	-0.420060	-0.145920	-0.21666	-0.035429	
32558	1.423610	-0.420060	-0.145920	-0.21666	-0.035429	
32559	-1.215643	-0.420060	-0.145920	-0.21666	-1.655225	
32560	0.983734	-0.420060	1.888424	-0.21666	-0.035429	

Workclass_0	Workclass_1	Workclass_2	Workclass_3	Workclass_4	...	\
-------------	-------------	-------------	-------------	-------------	-----	---

0	0	0	0	0	0	0 ...
1	0	0	0	0	0	0 ...
2	0	0	0	0	0	1 ...
3	0	0	0	0	0	1 ...
4	0	0	0	0	0	1 ...
...
32556	0	0	0	0	0	1 ...
32557	0	0	0	0	0	1 ...
32558	0	0	0	0	0	1 ...
32559	0	0	0	0	0	1 ...
32560	0	0	0	0	0	0 ...

	Country_32	Country_33	Country_34	Country_35	Country_36	Country_37	\
0	0	0	0	0	0	0	
1	0	0	0	0	0	0	
2	0	0	0	0	0	0	
3	0	0	0	0	0	0	
4	0	0	0	0	0	0	
...	
32556	0	0	0	0	0	0	
32557	0	0	0	0	0	0	
32558	0	0	0	0	0	0	
32559	0	0	0	0	0	0	
32560	0	0	0	0	0	0	

	Country_38	Country_39	Country_40	Country_41
0	0	1	0	0
1	0	1	0	0
2	0	1	0	0
3	0	1	0	0
4	0	0	0	0
...
32556	0	1	0	0
32557	0	1	0	0
32558	0	1	0	0
32559	0	1	0	0
32560	0	1	0	0

[32561 rows x 91 columns]

```
[9]: # One-hot encoding of categorical data
X = pd.get_dummies(X)

# Conversion of data frame to numpy
X = X.values

# Converision: {False, True} --> {0., 1.}
```

```
y = y.astype(float)
```

```
[10]: print(X.shape)
      print(y.shape)
      #print(X)
      #print(y)
      print(max(X[:,42]))
```

```
(32561, 91)
```

```
(32561,)
```

```
1.0
```

0.4 train-val-test split

```
[11]: from sklearn.model_selection import train_test_split

X_,X_test,y_,y_test = train_test_split(X,y,test_size=1/10,stratify=y)
#X_,X_test,y_,y_test = train_test_split(X,y,test_size=1/10)

X_train,X_val,y_train,y_val = train_test_split(X_,y_,test_size=1/9,stratify=y_)
#X_train,X_val,y_train,y_val = train_test_split(X_,y_,test_size=1/9)

print(X_train.shape)
print(X_val.shape)
print(X_test.shape)
print(sum(y_train)/y_train.shape)
print(sum(y_val)/y_val.shape)
print(sum(y_test)/y_test.shape)
```

```
(26048, 91)
```

```
(3256, 91)
```

```
(3257, 91)
```

```
[0.24082463]
```

```
[0.24078624]
```

```
[0.24071231]
```

0.5 Logistic regression

```
[12]: from sklearn.linear_model import LogisticRegression
```

```
[13]: model_LR = LogisticRegression()

# training
model_LR.fit(X_train, y_train)

# evaluation
val_acc = model_LR.score(X_val, y_val)
```



```
print(val_acc)
```

```
0.8507371007371007
```

```
lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
[14]: from joblib import dump  
  
      dump(model_LR, 'LR_sample.joblib')
```

```
[14]: ['LR_sample.joblib']
```

```
[15]: from joblib import load  
  
      loaded_model_LR = load('LR_sample.joblib')
```

```
[ ]:
```