

Convolutional neural networks

Practice Session 7

Changho Suh

January 24, 2024

Recap: Two building blocks of CNNs

1. Convolutional layer (Conv layer)

Role: Mimick neurons' behaviors:
Reacting only to receptive fields.

2. Pooling layer

Role: Downsample to reduce complexity
(# parameters & memory size).

Recap: Conv layer

1. A neuron reacts only to a **receptive field**.
2. A neighboring neuron concerns a receptive field shifted by **stride**.
3. Two types of **zero padding**: same, valid
4. Consists of a 2D **feature map**
5. Use of **multiple filters**
6. Has three **RGB** channels for colored images

Recap: Pooling layer

1. Two types of pooling:
Max pooling, average pooling
2. A default choice: 2×2 pooling filter, stride=2
3. Works on each channel independently


Recap: Two popular CNNs

1. **AlexNet** (2012)

Won the ImageNet competition.



Alex Krizhevsky



Ilya Sutskever




Geoffrey Hinton

Anchored the deep learning revolution.

2. **ResNet** (2015)

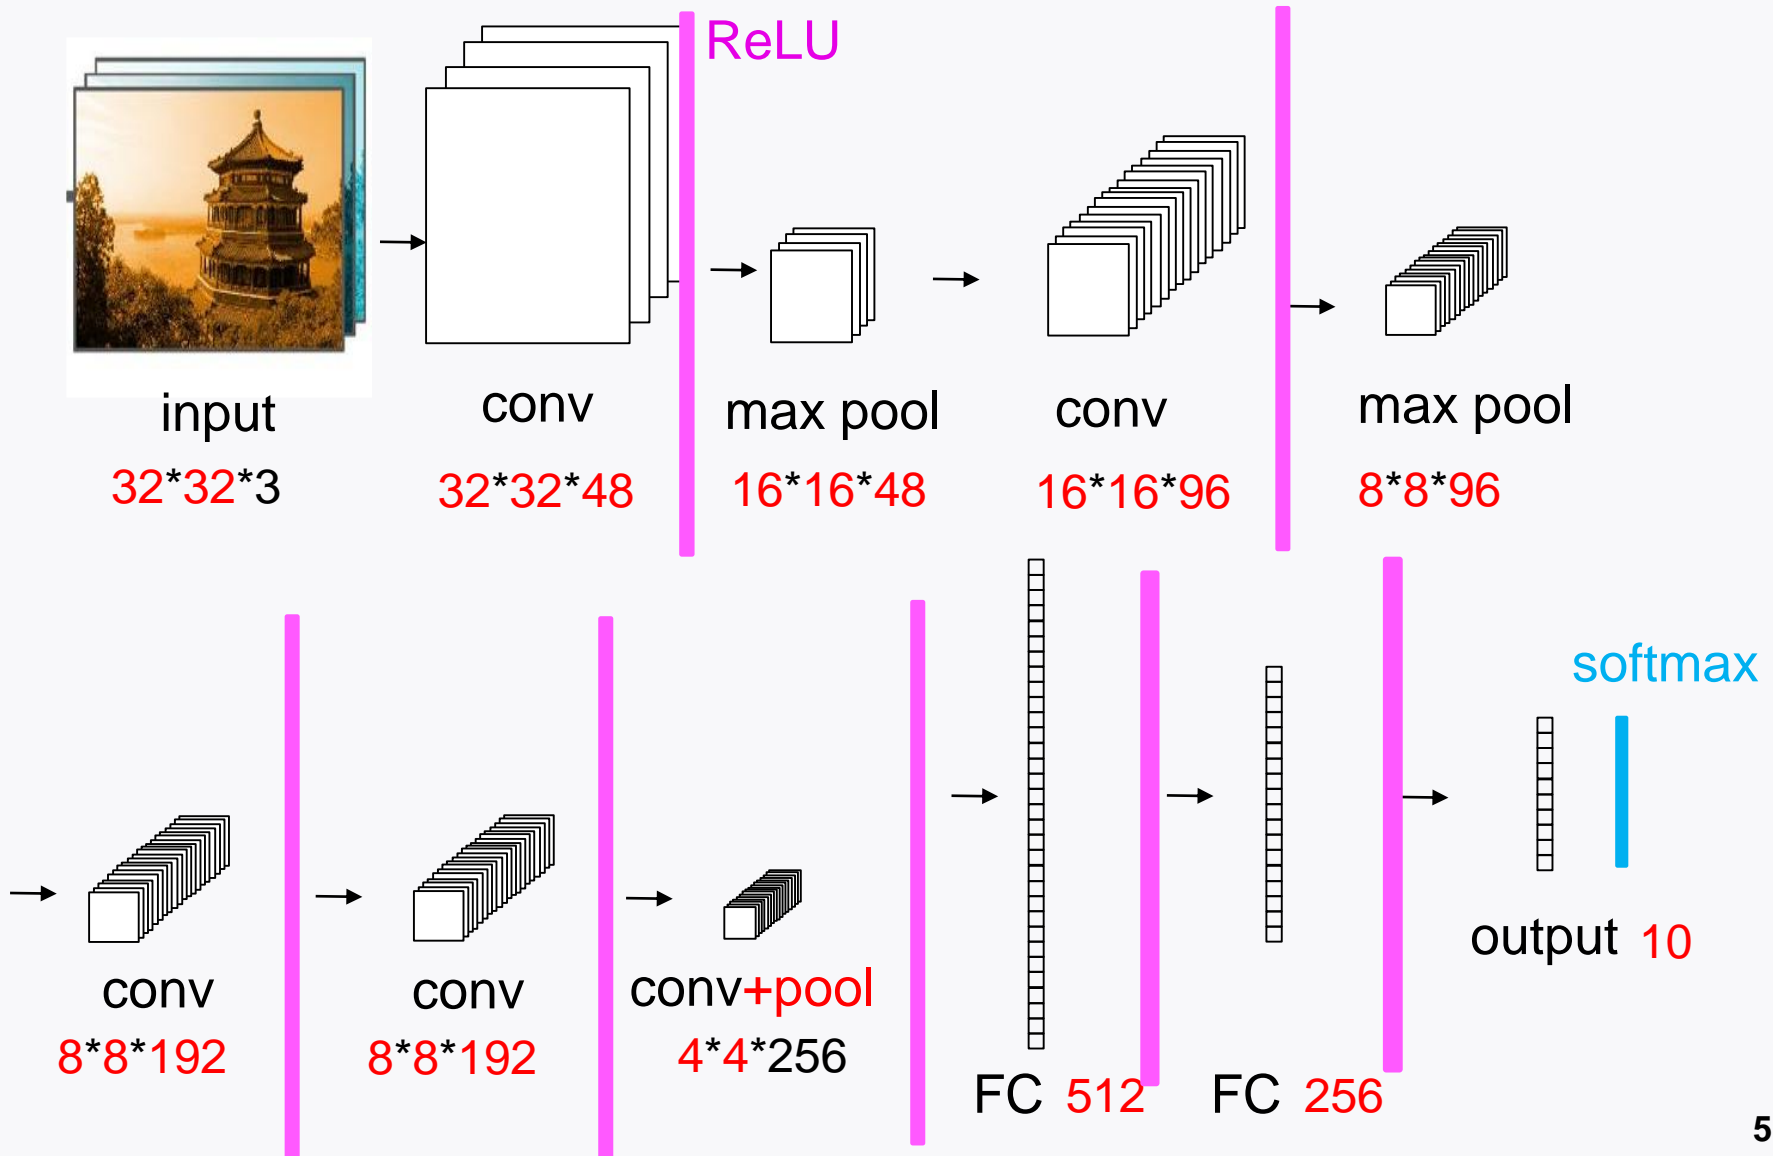
Won the 2015 ImageNet competition.

Currently the most powerful & arguably the simplest!

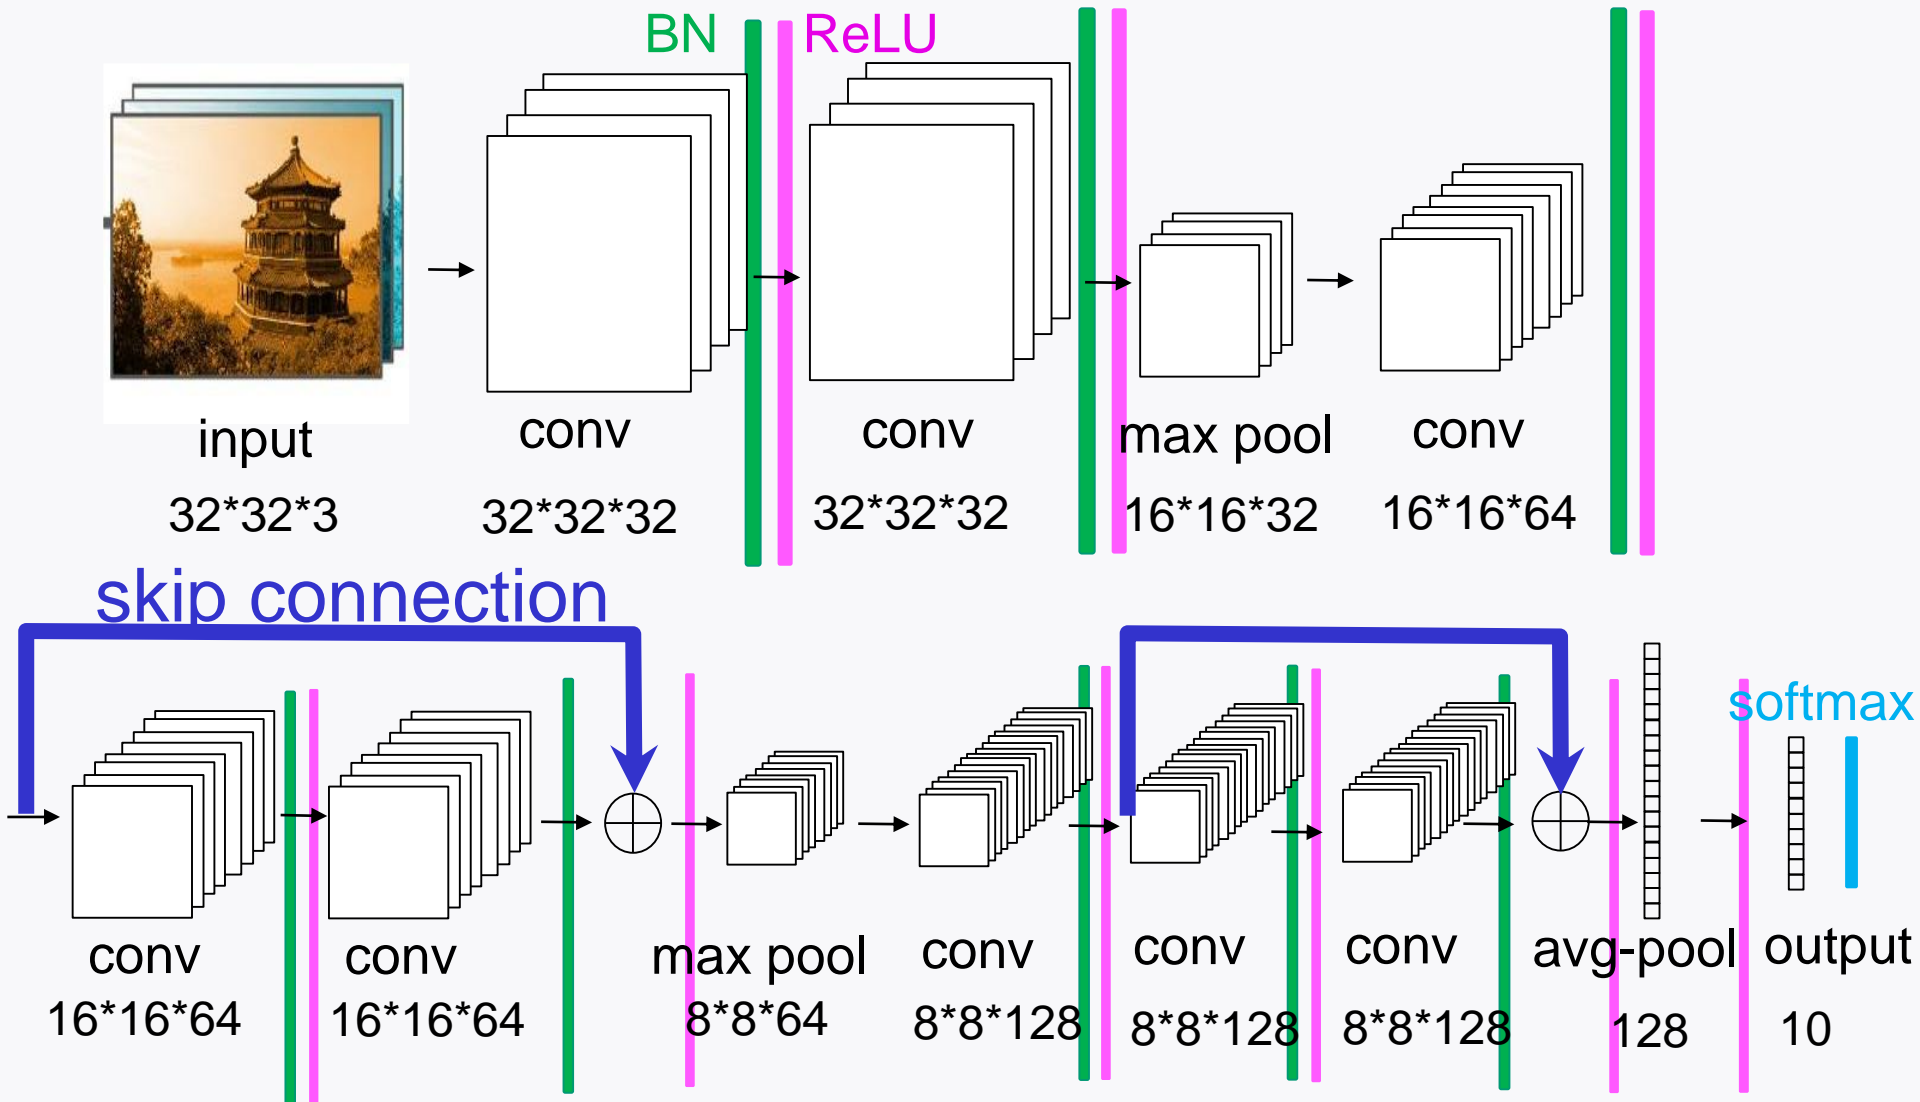


Kaiming He

Recap: Simplified AlexNet



Recap: Simplified ResNet



Outline

Will implement three prominent models:

1. LeNet5

Task: Handwritten digit classification (MNIST)


2. Simplified AlexNet

Task: Image recognition (CIFAR10)

3. Simplified ResNet

Task: Image recognition (CIFAR10)


Handwritten digit classification



MNIST dataset


$$\{(x^{(i)}, y^{(i)})\}_{i=1}^m \quad m = 60,000$$

Examples:



Yann LeCun 1998


Input image & label



How to load MNIST dataset

```
from tensorflow.keras.datasets import mnist  
  
(X_train, y_train), (X_test, y_test) = mnist.load_data()  
  
X_train, X_test = X_train/255., X_test/255
```

LeNet-5: CNN model



LeNet-5: Tensorflow coding


```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Flatten, Dense
from tensorflow.keras.layers import Conv2D
from tensorflow.keras.layers import MaxPool2D
```

```
model_lenet = Sequential()
```

```
#1st stack ([Conv]+[ReLU]+[Pool])
```

```
model_lenet.add(Conv2D(input_shape=(28,28,1),
                        kernel_size=(5,5),
                        strides=(1,1),
                        filters=32,
                        padding='same',
                        activation='relu'
                    ))
```

```
model_lenet.add(MaxPool2D(pool_size=(2,2),
                           strides=(2,2),
                           padding='valid'))
```




LeNet-5: Tensorflow coding

#2nd stack ([Conv]+[ReLU]+[Pool])

```
model_lenet.add(Conv2D(kernel_size=(5,5),
                        strides=(1,1),
                        filters=48,
                        padding='same',
                        activation='relu'
                        ))

model_lenet.add(MaxPool2D(pool_size=(2,2),
                           strides=(2,2),
                           padding='valid'))
```



Three fully Connected Layers

```
model_lenet.add(Flatten())
model_lenet.add(Dense(256, activation='relu' ))
model_lenet.add(Dense(84, activation='relu'))
model_lenet.add(Dense(10, activation='softmax'))
```

Compile

```
from tensorflow.keras.optimizers import Adam

opt = Adam(learning_rate = 0.001,
           beta_1 = 0.9,
           beta_2 = 0.999)

model_lenet.compile(optimizer = opt,
                    loss='sparse_categorical_crossentropy',
                    metrics=['acc'])
```


Training & evaluation

```
# Conversion from 3D to 4D tensor
X_train, X_test = X_train.reshape(-1,28,28,1), X_test.reshape(-1,28,28,1)

# Training
model_lenet.fit(X_train,y_train,epochs=10)

# Evaluation
test_performance = model_lenet.evaluate(X_test, y_test)
print(test_performance)
```

```
[0.036696918308734894, 0.9922999739646912]
```

Look ahead

Will implement simplified AlexNet in the context of image recognition (CIFAR10).