

# Convolutional neural networks

## Practice Session 9

Changho Suh

January 24, 2024


# Outline

---

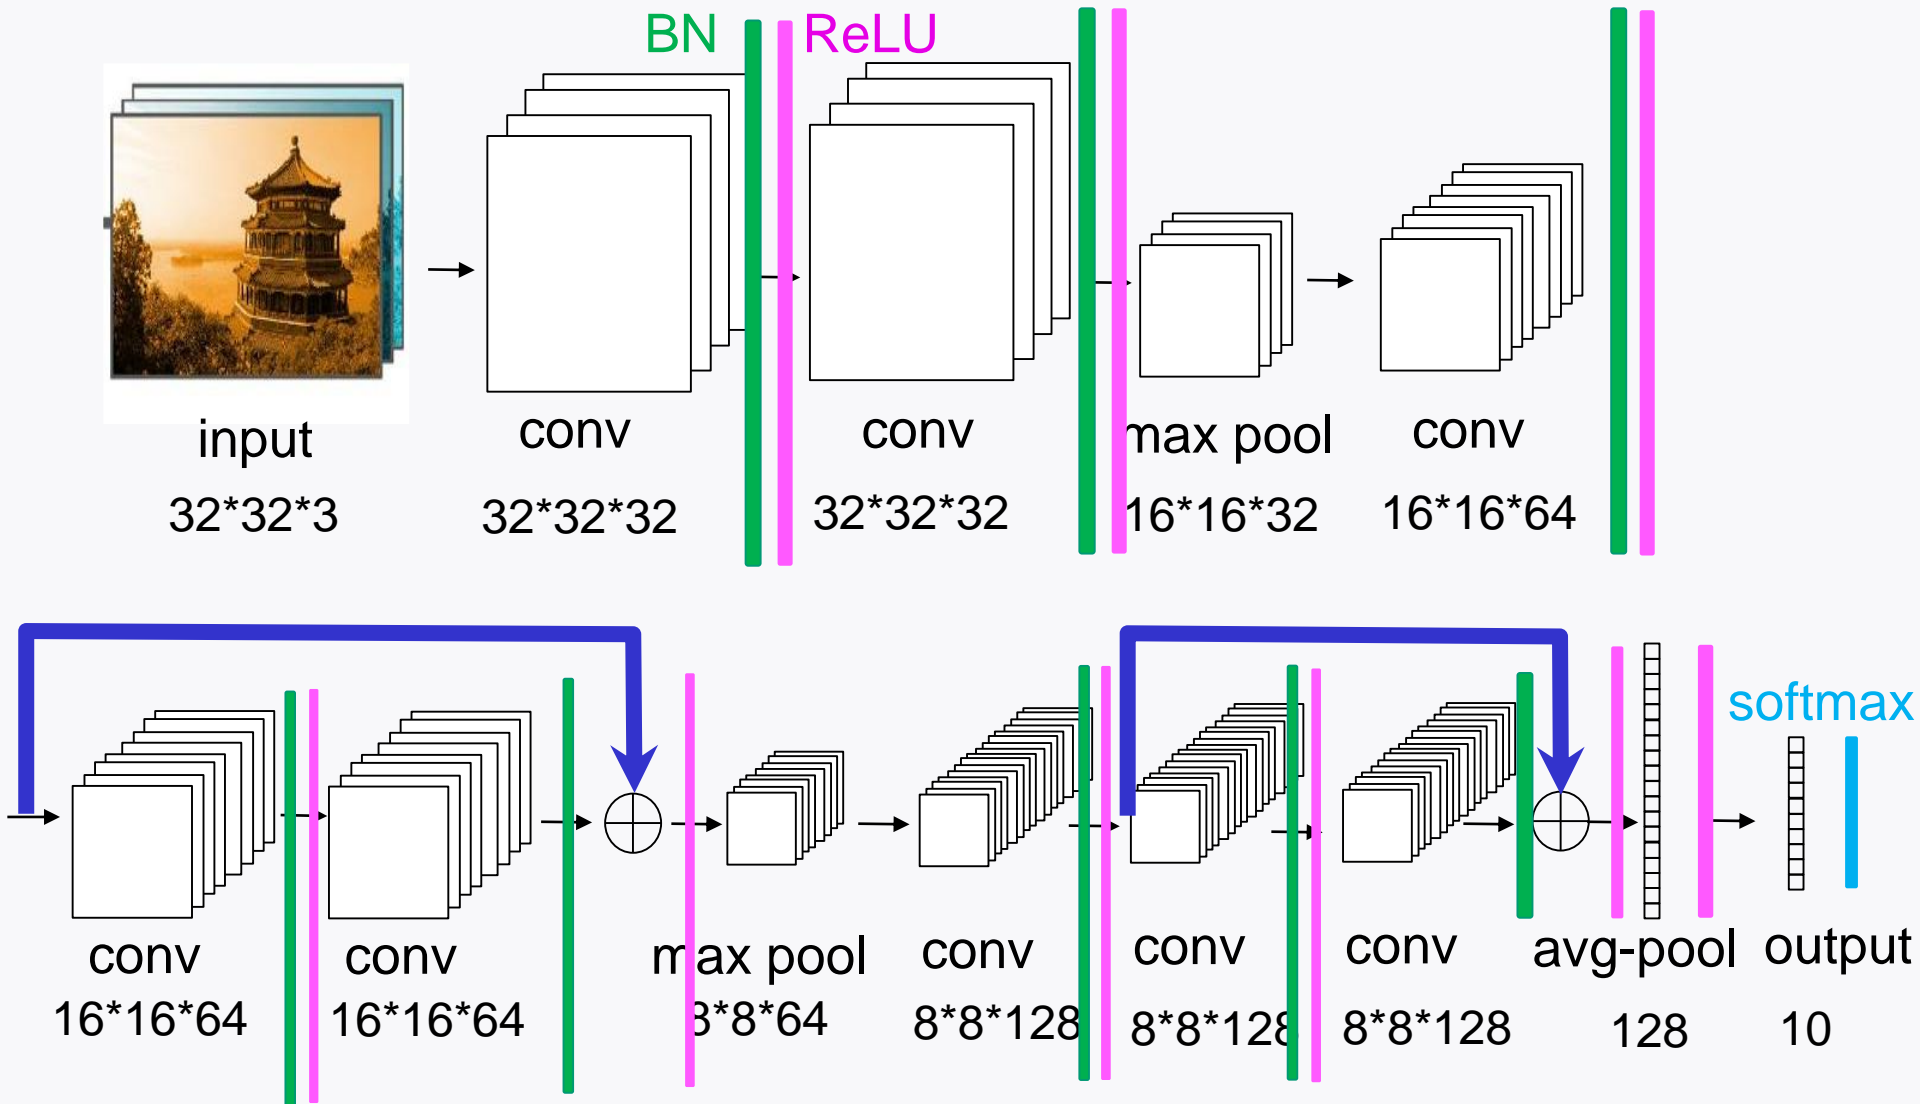
Will implement the simplified ResNet.

Task: Image recognition (CIFAR10)

# Image recognition



# Simplified ResNet



# Two ways to construct a model

## 1. Sequential:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Flatten
from tensorflow.keras.layers import Dense

Model_NN = Sequential()

Model_NN.add(Flatten(input_shape=(28,28)))

Model_NN.add(Dense(128, activation='relu'))

Model_NN.add(Dense(10, activation='softmax'))
```

# Two ways to construct a model

## 2. Model:

```
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Flatten
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import Input

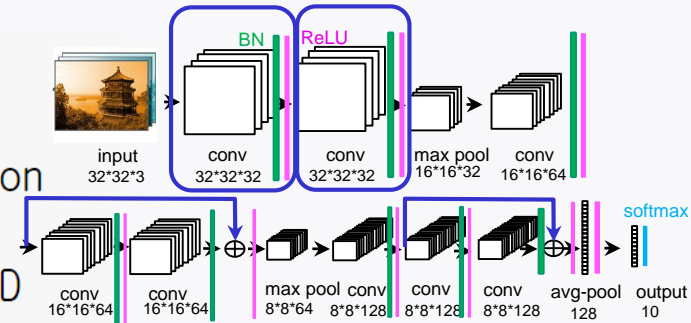
inputs = Input(shape=(28,28))
flatten = Flatten()(inputs)
hidden = Dense(128, activation='relu')(flatten)
outputs = Dense(10, activation='softmax')(hidden)
model_NN = Model(inputs = inputs, outputs = outputs)
```

# ResNet: Tensorflow coding

```

from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input
from tensorflow.keras.layers import BatchNormalization
from tensorflow.keras.layers import ReLU
from tensorflow.keras.layers import Conv2D, MaxPool2D
from tensorflow.keras.layers import Add, AveragePooling2D

```



```

inputs = Input(shape=(32,32,3))
x = Conv2D(32, kernel_size=(3,3), strides=(1,1),
          padding='same', use_bias=False)(inputs)
x = BatchNormalization()(x)
x = ReLU()(x)

x = Conv2D(32, kernel_size=(3,3), strides=(1,1),
          padding='same', use_bias=False)(x)
x = BatchNormalization()(x)
x = ReLU()(x)

x = Conv2D(64, kernel_size=(3,3), strides=(1,1),
          padding='same', use_bias=False)(x)
x = BatchNormalization()(x)
x = ReLU()(x)

x = Conv2D(64, kernel_size=(3,3), strides=(1,1),
          padding='same', use_bias=False)(x)
x = BatchNormalization()(x)
x = ReLU()(x)

x = Conv2D(128, kernel_size=(3,3), strides=(1,1),
          padding='same', use_bias=False)(x)
x = BatchNormalization()(x)
x = ReLU()(x)

x = Conv2D(128, kernel_size=(3,3), strides=(1,1),
          padding='same', use_bias=False)(x)
x = BatchNormalization()(x)
x = ReLU()(x)

x = AveragePooling2D()(x)
x = Softmax()(x)

```

```
x = Add()(x, skip)
x = ReLU()(x)
```





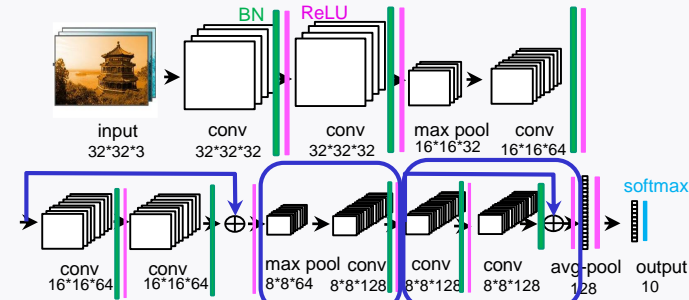
# ResNet: Tensorflow coding

```
x = MaxPool2D(pool_size=(2,2), strides=(2,2),
              padding='valid')(x)
x = Conv2D(128, kernel_size=(3,3), strides=(1,1),
          padding='same', use_bias=False)(x)
x = BatchNormalization()(x)
x = ReLU()(x)
```

*# 2nd skip connection*

```
skip = x
x = Conv2D(128, kernel_size=(3,3), strides=(1,1),
          padding='same', use_bias=False)(x)
x = BatchNormalization()(x)
x = ReLU()(x)
x = Conv2D(128, kernel_size=(3,3), strides=(1,1),
          padding='same', use_bias=False)(x)
x = BatchNormalization()(x)


x = Add()(x, skip)
x = ReLU()(x)
```



# ResNet: Tensorflow coding

*# Average Pooling*

```
x = AveragePooling2D(pool_size=(8,8))(x)
x = Flatten()(x)
x = ReLU()(x)
outputs = Dense(10, activation='softmax')(x)
```



```
Model_resnet = Model(inputs = inputs, outputs = outputs)
```

# Compile

```
from tensorflow.keras.optimizers import Adam

opt = Adam(learning_rate = 0.001,
           beta_1 = 0.9,
           beta_2 = 0.999)

Model_resnet.compile(optimizer = opt,
                     loss='sparse_categorical_crossentropy',
                     metrics=['acc'])
```

# Training & evaluation

```
# training
```

```
Model_resnet.fit(X_train,y_train,epochs=20)
```

```
# Evaluation
```

```
test_performance = Model_resnet.evaluate(X_test,y_test)  
print(test_performance)
```

```
[0.7550032138824463, 0.8246999979019165]
```