

PS15

January 24, 2023

0.1 Random forests

0.2 Iris plants classification

```
[1]: from sklearn.datasets import load_iris
iris = load_iris()
X = iris.data
y = iris.target
```

```
[2]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
#from sklearn.metrics import accuracy_score
```

```
[7]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2)

rnd_clf = RandomForestClassifier(n_estimators=500, max_leaf_nodes=16, n_jobs=-1)
# n_jobs=-1 : use all available CPU's
rnd_clf.fit(X_train, y_train)

y_pred = rnd_clf.predict(X_test)
print(rnd_clf.score(X_test,y_test))

#print(accuracy_score(y_pred, y_test))
```

1.0

1.0

```
[8]: for name, score in zip(iris.feature_names, rnd_clf.feature_importances_):
    print(name, score)
```


sepal length (cm) 0.10810853310073656
sepal width (cm) 0.025670524096247496
petal length (cm) 0.44972571693418767
petal width (cm) 0.4164952258688282

```
[12]: import matplotlib.pyplot as plt
import numpy as np
```

```
sorted_idx = rnd_clf.feature_importances_.argsort()
plt.figure(figsize=(5,3), dpi=150)
plt.barh(np.asarray(iris.feature_names)[sorted_idx],
         rnd_clf.feature_importances_[sorted_idx])

#sorted_idx = rnd_clf.feature_importances_.argsort()[::-1]
#plt.barh(np.asarray(iris.feature_names)[sorted_idx][::-1],
#         rnd_clf.feature_importances_[sorted_idx][::-1])
```

[12]: <BarContainer object of 4 artists>



0.3 California Housing Price prediction

```
[13]: from sklearn.datasets import fetch_california_housing
cali_prices = fetch_california_housing()
X_reg = cali_prices.data
y_reg = cali_prices.target
```

```
[14]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
```

```
[15]: X_train,X_test,y_train,y_test = train_test_split(X_reg,y_reg,test_size = 0.01)
for_reg = RandomForestRegressor()
for_reg.fit(X_train, y_train)
y_pred_train = for_reg.predict(X_train)
y_pred_test = for_reg.predict(X_test)
print(mean_squared_error(y_pred_train, y_train))
```

```
print(mean_squared_error(y_pred_test, y_test))
```

0.03426737442063849

0.21248308755938008

0.4 MNIST classification

```
[18]: from keras.datasets import mnist
      from sklearn.model_selection import train_test_split

      (X_train, y_train), (X_test, y_test) = mnist.load_data()
      X_train, X_test = X_train / 255., X_test / 255.

      X_train = X_train.reshape(-1, 28*28)
      X_test = X_test.reshape(-1, 28*28)

      #X_train = X_train.reshape(X_train.shape[0], -1)
      #X_test = X_test.reshape(X_test.shape[0], -1)
```

```
[20]: from sklearn.ensemble import RandomForestClassifier
      from sklearn.metrics import accuracy_score


      rnd_clf = RandomForestClassifier(n_estimators=500,
                                     max_leaf_nodes=16,
                                     n_jobs=-1)

      rnd_clf.fit(X_train, y_train)
      y_pred = rnd_clf.predict(X_test)
      print(rnd_clf.score(X_test, y_test))
      #print(accuracy_score(y_pred, y_test))
```

0.8309

```
[16]: import matplotlib.pyplot as plt

      feature_importances = rnd_clf.feature_importances_
      feature_importances = feature_importances.reshape(28, 28)
      plt.imshow(feature_importances, cmap='hot')
      plt.colorbar(label='importance')
      plt.show()
```



0.5 Hyperparameter search: GridSearchCV

```
[17]: from sklearn.datasets import load_iris
      from sklearn.model_selection import train_test_split
```

```
Iris = load_iris()
X = Iris.data
y = Iris.target
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2)
```

```
[18]: from sklearn.model_selection import GridSearchCV
```

```
[19]: param_grid = {'n_estimators': [3,10,100,500],
                  'max_features': [0.25, 0.5, 0.75, 1]}

forest_clf = RandomForestClassifier(max_depth=2)
grid_search = GridSearchCV(forest_clf,param_grid,cv=5,scoring='accuracy')
grid_search.fit(X_train, y_train)
```

```
[19]: GridSearchCV(cv=5, estimator=RandomForestClassifier(max_depth=2),
                  param_grid={'max_features': [0.25, 0.5, 0.75, 1],
                              'n_estimators': [3, 10, 100, 500]},
                  scoring='accuracy')
```

```
[20]: print(grid_search.best_params_)
```

```
{'max_features': 0.25, 'n_estimators': 100}
```

```
[21]: print(grid_search.best_estimator_)
```

```
RandomForestClassifier(max_depth=2, max_features=0.25)
```

```
[22]: print(grid_search.best_score_)
```

```
0.9416666666666667
```

```
[23]: feature_importances = grid_search.best_estimator_.feature_importances_  
print(feature_importances)
```

```
[0.21406567 0.15593039 0.32120427 0.30879968]
```

```
[24]: y_pred = grid_search.best_estimator_.predict(X_test)  
print(accuracy_score(y_pred, y_test))
```

```
0.9666666666666667
```

0.6 Hyperparameter search: RandomizedSearchCV

```
[25]: from sklearn.datasets import load_iris  
from sklearn.model_selection import train_test_split  
Iris = load_iris()  
X = Iris.data  
y = Iris.target  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
[26]: from sklearn.model_selection import RandomizedSearchCV
```

```
[27]: param_distributions = {'n_estimators': range(500),  
                           'max_features': range(1, 5)}  
  
forest_clf = RandomForestClassifier(max_depth=2)  
  
randomized_search = RandomizedSearchCV(forest_clf,  
                                       param_distributions,  
                                       cv=4, n_iter = 50,  
                                       scoring='accuracy')  
  
randomized_search.fit(X_train, y_train)
```

```
[27]: RandomizedSearchCV(cv=4, estimator=RandomForestClassifier(max_depth=2),  
                       n_iter=50,  
                       param_distributions={'max_features': range(1, 5),  
                                           'n_estimators': range(0, 500)},  
                       scoring='accuracy')
```

```
[28]: print(randomized_search.best_params_)
```

```
{'n_estimators': 212, 'max_features': 2}
```

```
[29]: print(randomized_search.best_estimator_)
```

```
RandomForestClassifier(max_depth=2, max_features=2, n_estimators=212)
```

```
[30]: print(randomized_search.best_score_)
```

```
0.9750000000000001
```

```
[31]: feature_importances = randomized_search.best_estimator_.feature_importances_  
print(feature_importances)
```

```
[0.1202956  0.00995066 0.43334277 0.43641097]
```

```
[32]: y_pred = randomized_search.best_estimator_.predict(X_test)
```

```
print(accuracy_score(y_pred, y_test))
```

```
0.9333333333333333
```