

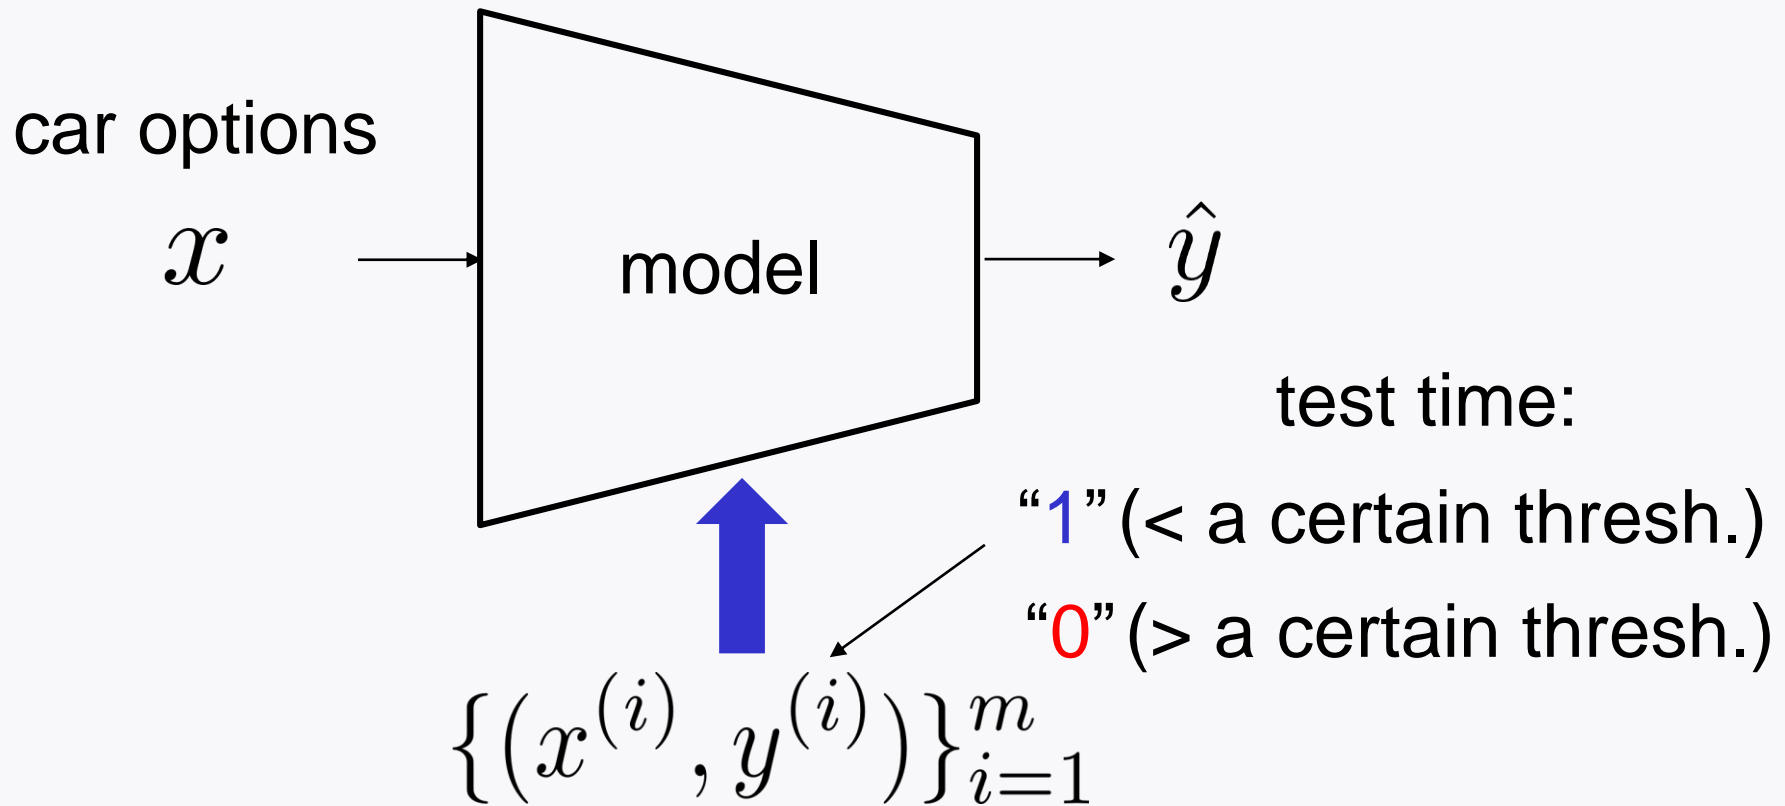
# Mini-project #1

## Practice Session 18

Changho Suh

January 29, 2024

# Recap: Test-time prediction



# Recap: Loading MB dataset

```
import pandas as pd
data = pd.read_csv('mercedes_test.csv')
data
```

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377	X378	X379	X380	X382	X383	X384	X385
0	0	130.81	k	v	at	a	d	u	j	o	...	0	0	1	0	0	0	0	0	0	0
1	6	88.53	k	t	av	e	d	y	l	o	...	1	0	0	0	0	0	0	0	0	0
2	7	76.26	az	w	n	c	d	x	j	x	...	0	0	0	0	0	0	1	0	0	0
3	9	80.62	az	t	n	f	d	x	l	e	...	0	0	0	0	0	0	0	0	0	0
4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0	0	0	0	0	0	0	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
4204	8405	107.39	ak	s	as	c	d	aa	d	q	...	1	0	0	0	0	0	0	0	0	0
4205	8406	108.77	j	o	t	d	d	aa	h	h	...	0	1	0	0	0	0	0	0	0	0
4206	8412	109.22	ak	v	r	a	d	aa	g	e	...	0	0	1	0	0	0	0	0	0	0
4207	8415	87.48	al	r	e	f	d	aa	l	u	...	0	0	0	0	0	0	0	0	0	0
4208	8417	110.85	z	r	ae	c	d	aa	g	w	...	1	0	0	0	0	0	0	0	0	0

4209 rows × 378 columns

strings (categorical data)

$$m = 4209 \quad n = 376 (= 378 - 2)$$

# Recap: Preprocessing

```
# Choose categorical data columns
cf = data.select_dtypes(include=['object']).columns
# To change it into "categorical" data type
data[cf]=data[cf].astype('category')
# One hot encoding
data = pd.get_dummies(data)
# Obtain X from data (excluding 'ID' and 'y')
X_df = data.drop(['ID', 'y'],axis=1)
# Obtain y from data
y_df = data['y']

# Convert y_df into binary labels
import numpy as np
TF_vector= (y_df<np.median(y_df))
y_df=TF_vector.astype(float)

# Conver data frame into numpy array
X,y = X_df.values, y_df.values
```

# Recap: Split into train and test datasets

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.1,stratify=y)
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(3788, 563)
(421, 563)
(3788,)
(421,)
```

# Designing an DNN model

---

Many hyperparameters to search.

Difficult to develop a well-performing model.

**Hence:** Important to do very *quick* and *convenient* search for hyperparameters.

There is a tool that may be useful for this:

**TensorBoard**

# TensorBoard



A **visualization** tool for  
Keras models

Three key features:

1. Draws curves with logs info
2. Provides visualization for models
3. Comes integrated with TensorFlow

# How to use TensorBoard?

1. Load TensorBoard notebook extension:

```
%load_ext tensorboard
```

2. Train a Keras model using a command **TensorBoard**.

3. Launch a TensorBoard browser:

```
%tensorboard --logdir=logs/
```

← directory where logs info will be saved



# TensorBoard in action: Load TensorBoard

---

```
%load_ext tensorboard
```

# TensorBoard in action: Train a Keras model

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.regularizers import l2
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.callbacks import LearningRateScheduler

# Construction of an DNN model
def create_model(n_layer=2, lambda_=0):
    model = Sequential()
    for i in range(n_layer-1):
        model.add(Dense(10, activation='relu',
                        kernel_regularizer=l2(lambda_), bias_regularizer=l2(lambda_)))

    model.add(Dense(1, activation='sigmoid',
                    kernel_regularizer=l2(lambda_), bias_regularizer=l2(lambda_)))
    return model
```

# TensorBoard in action: Train a Keras model

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.regularizers import l2
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.callbacks import LearningRateScheduler
import os
from tensorflow.keras.callbacks import TensorBoard

model = create_model(n_layer=2, lambda_=0)
opt = Adam(learning_rate=1e-3)
model.compile(optimizer=opt,
              loss='binary_crossentropy',
              metrics=['acc'])

# generate a path directory where logs info will be saved
logdir = os.path.join('logs', '[1]no_regularization')
tb_callback = TensorBoard(logdir)
```

To save logs info in the directory “logdir”

# TensorBoard in action: Train a Keras model

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.regularizers import l2
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.callbacks import LearningRateScheduler
import os
from tensorflow.keras.callbacks import TensorBoard

model = create_model(n_layer=2, lambda_=0)
opt = Adam(learning_rate=1e-3)
model.compile(optimizer=opt,
              loss='binary_crossentropy',
              metrics=['acc'])

# generate a path directory where logs info will be saved
logdir = os.path.join('logs', '[1]no_regularization') logs\[1]no_regularization
tb_callback = TensorBoard(logdir)
es_callback = EarlyStopping(monitor='val_acc', patience=20)
hist = model.fit(X_train, y_train,
                 validation_split=1/9, epochs=100,
                 verbose=0, callbacks=[tb_callback, es_callback])
model.evaluate(X_test, y_test)
```

# TensorBoard in action: Launch a TB browser

```
%tensorboard --logdir=logs/
```

TensorBoard interface showing the SCALARS tab. The interface includes a top navigation bar with tabs for TensorBoard, SCALARS (selected), GRAPHS, TIME SERIES, and INACTIVE. A large text overlay "click this tap" points to the SCALARS tab.

On the left sidebar, there are settings for "Show data download links" (unchecked), "Ignore outliers in chart scaling" (checked), "Tooltip sorting method" (default), "Smoothing" (0.599), and "Horizontal Axis" (STEP selected). Below these are "Runs" and "WALL" sections.

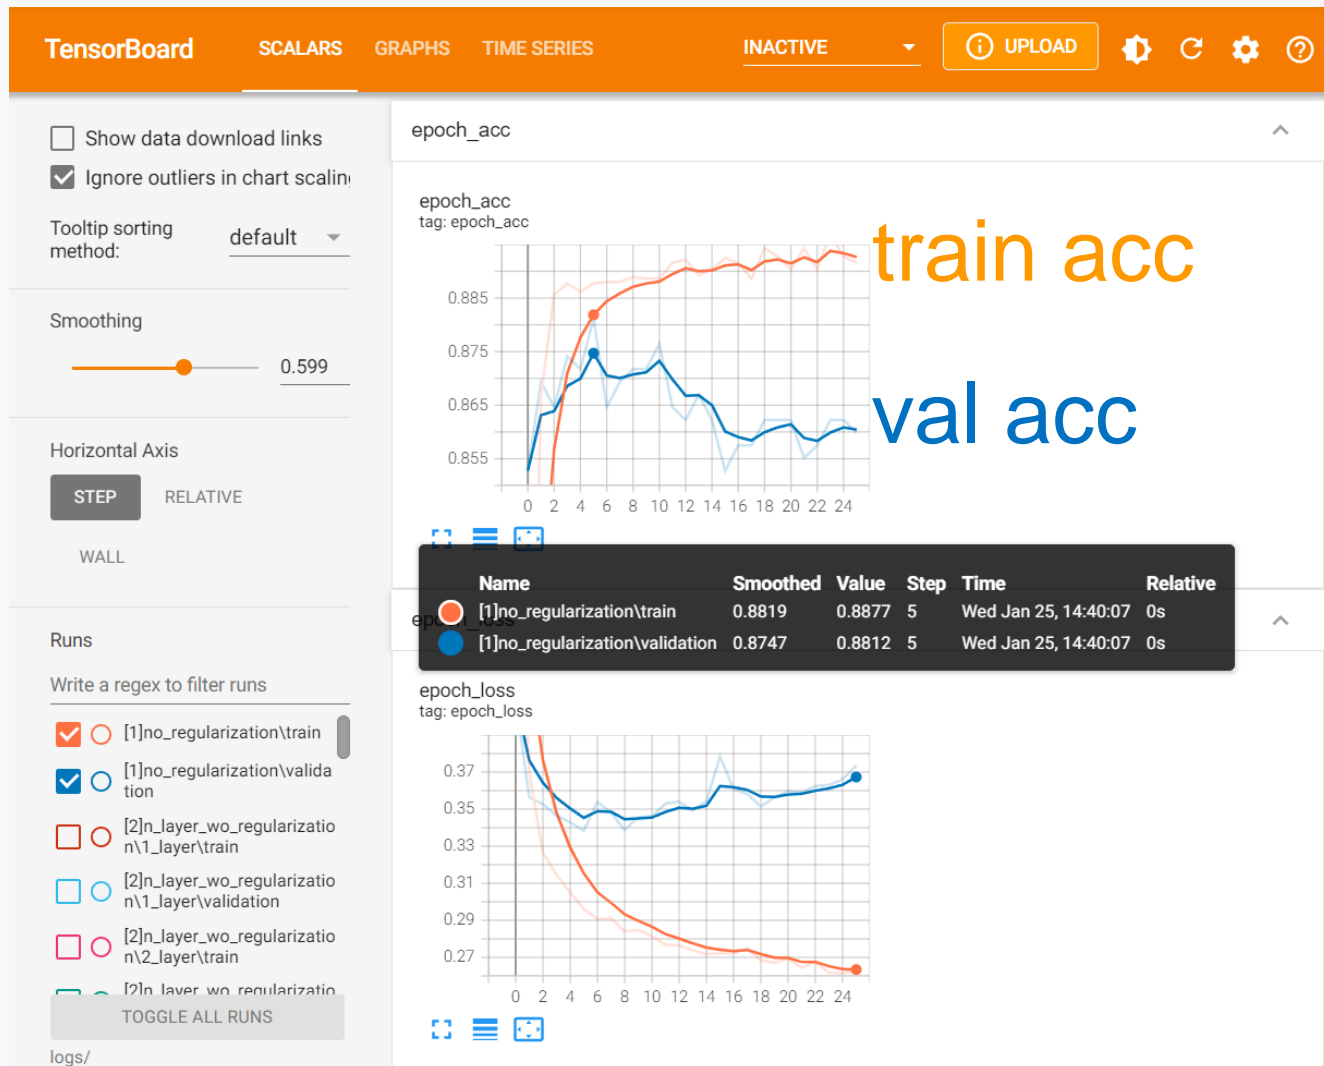
The main display area shows two line charts. The top chart is titled "epoch\_acc" and the bottom chart is titled "epoch\_loss". Both charts show training and validation metrics over 24 epochs. The "epoch\_acc" chart shows training accuracy (blue line) and validation accuracy (orange line). The "epoch\_loss" chart shows training loss (blue line) and validation loss (orange line).

Below the charts, there is a table of runs. The first run is highlighted:

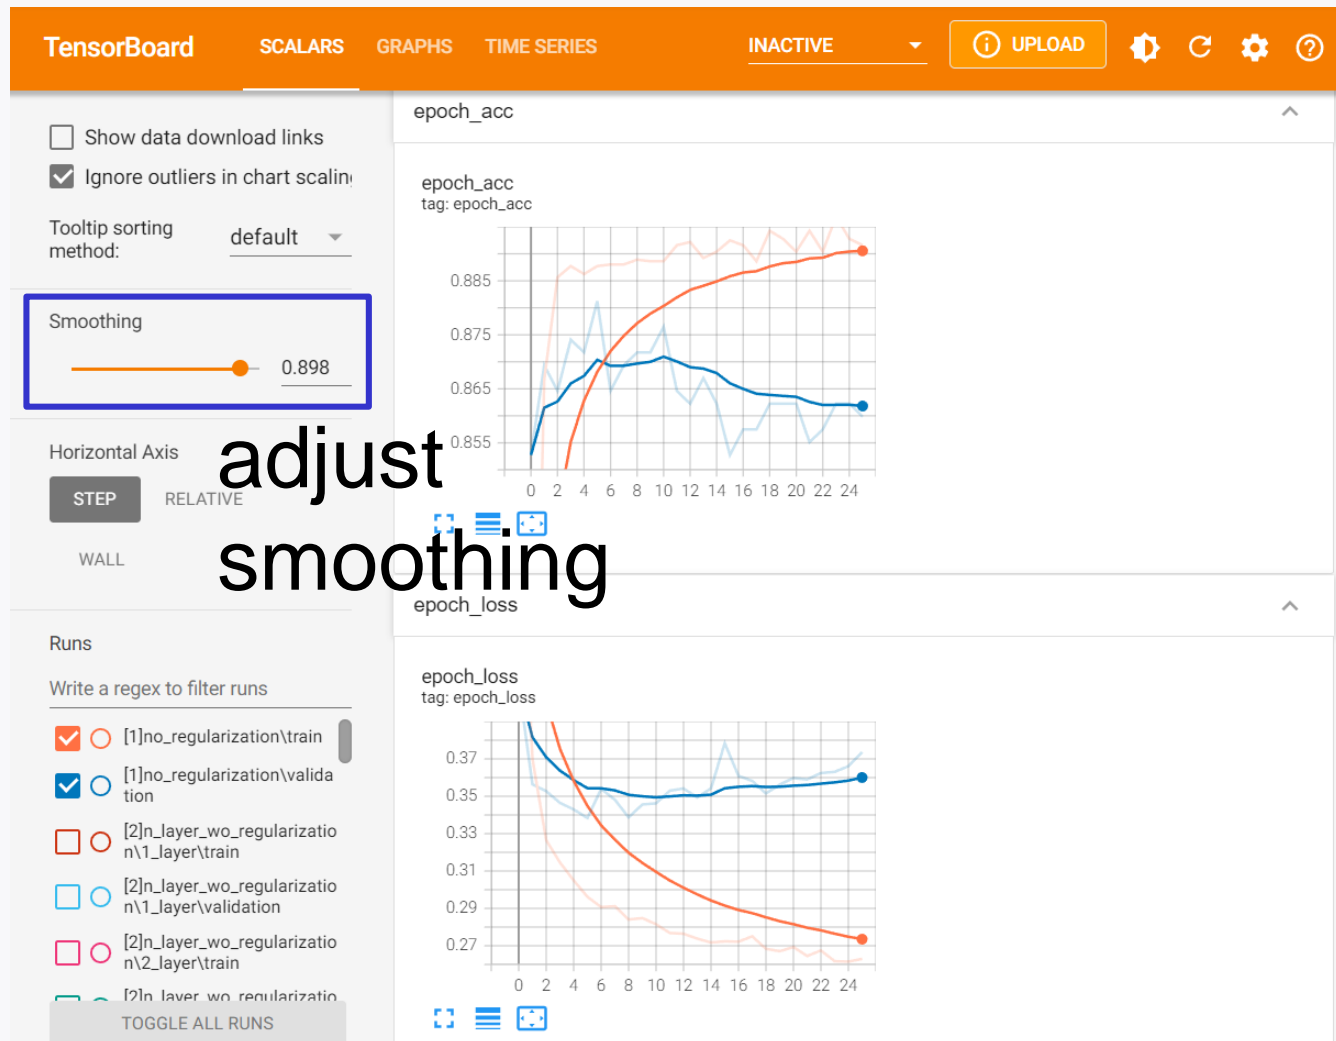
이름	수정한 날짜	유형
[1]no_regularization	2023-01-25 오후 2:40	파일 폴더

An arrow points from the "epoch\_loss" chart to the first run in the table.

# TensorBoard in action: Launch a TB browser



# TensorBoard in action: Launch a TB browser



# Hyperparameter search: Number of layers

```
n_layer_list = [1,2,3,4,5]

for n_layer in n_layer_list:
    model = create_model(n_layer=n_layer, lambda_=0)
    opt = Adam(learning_rate=1e-3)
    model.compile(optimizer=opt,
                  loss='binary_crossentropy',
                  metrics=['acc'])
    logdir = os.path.join('logs', '[2]n_layer_wo_regularization', '{}_layer'.format(n_layer))
    tb_callback = TensorBoard(logdir)
    es_callback = EarlyStopping(monitor='val_acc', patience=20)
    hist = model.fit(X_train, y_train,
                    validation_split=1/9, epochs = 100,
                    verbose=0, callbacks=[tb_callback, es_callback])
    model.evaluate(X_test, y_test)
```

```
14/14 [=====] - 0s 690us/step - loss: 0.3235 - acc: 0.8789
14/14 [=====] - 0s 921us/step - loss: 0.3347 - acc: 0.8694
14/14 [=====] - 0s 691us/step - loss: 0.3521 - acc: 0.8599
14/14 [=====] - 0s 691us/step - loss: 0.3421 - acc: 0.8765
14/14 [=====] - 0s 690us/step - loss: 0.3806 - acc: 0.8599
```



# Hyperparameter search: Number of layers

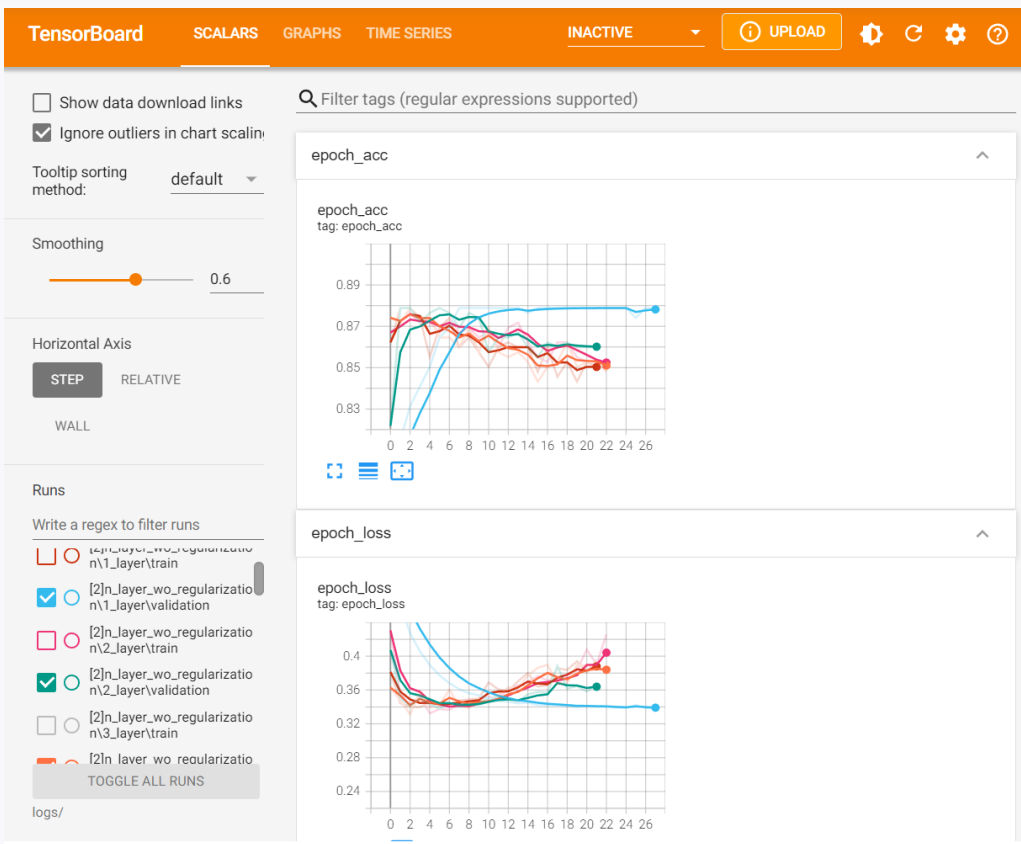
[현대자동차] > DS30\_5차수 > DAY6 > logs

이름	수정한 날짜	유형
[1]no_regularization	2023-01-25 오후 2:40	파일 폴더
[2]n_layer_wo_regularization	2023-01-25 오후 2:40	파일 폴더

[현대자동차] > DS30\_5차수 > DAY6 > logs > [2]n\_layer\_wo\_regularization >


이름	수정한 날짜	유형
1_layer	2023-01-25 오후 2:40	파일 폴더
2_layer	2023-01-25 오후 2:40	파일 폴더
3_layer	2023-01-25 오후 2:40	파일 폴더
4_layer	2023-01-25 오후 2:40	파일 폴더
5_layer	2023-01-25 오후 2:40	파일 폴더



```
%tensorboard --logdir=logs/[2]n_layer_wo_regularization
```



# Hyperparameter search: Number of layers

epoch\_acc  
tag: epoch\_acc



Name	Smoothed	Value
 [2]n_layer_wo_regularization\1_layer\validation	0.849	0.8646
 [2]n_layer_wo_regularization\2_layer\validation	0.8753	0.8789
 [2]n_layer_wo_regularization\3_layer\validation	0.8701	0.8646
 [2]n_layer_wo_regularization\4_layer\validation	0.8676	0.8694
 [2]n_layer_wo_regularization\5_layer\validation	0.87	0.867

# Hyperparameter search: Regularization factor

```
lambda_list = [1e-3, 1e-2, 1e-1, 1, 10]
for lambda_ in lambda_list:
    model = create_model(n_layer=2, lambda_=lambda_)
    opt = Adam(learning_rate=1e-3)
    model.compile(optimizer=opt,
                  loss='binary_crossentropy',
                  metrics=['acc'])
    logdir = os.path.join('logs', '[3]2_layer_w_regularization', 'lambda_{}'.format(lambda_))
    tb_callback = TensorBoard(logdir)
    es_callback = EarlyStopping(monitor='val_acc', patience=20)
    hist = model.fit(X_train, y_train,
                    validation_split=1/9, epochs = 100,
                    verbose=0, callbacks=[tb_callback, es_callback] )
    model.evaluate(X_test, y_test)
```

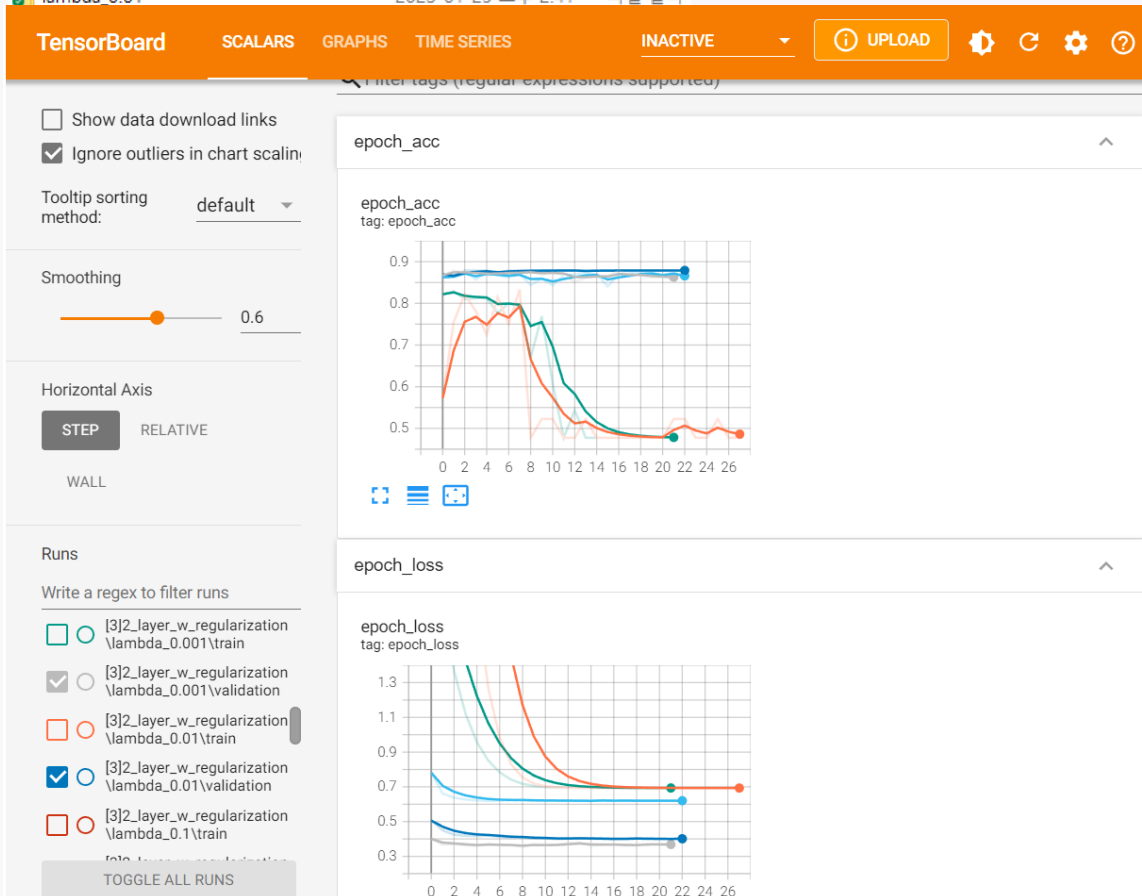
```
14/14 [=====] - 0s 614us/step - loss: 0.3545 - acc: 0.8622
14/14 [=====] - 0s 618us/step - loss: 0.3986 - acc: 0.8741
14/14 [=====] - 0s 691us/step - loss: 0.6349 - acc: 0.8599
14/14 [=====] - 0s 764us/step - loss: 0.6931 - acc: 0.5012
14/14 [=====] - 0s 691us/step - loss: 0.6932 - acc: 0.5012
```

# Hyperparameter search: Regularization factor

```
%tensorboard --logdir=logs/[3]2_layer_w_regularization
```

[현대자동차] > DS30\_5차수 > DAY6 > logs > [3]2\_layer\_w\_regularization >

이름	수정한 날짜	유형
lambda_0.001	2023-01-25 오후 2:41	파일 폴더
lambda_0.01	2023-01-25 오후 2:41	파일 폴더



# Hyperparameter search: Regularization factor



# Hyperparameter search: Learning rate

```
lr_list = [1e-1, 1e-2, 1e-3, 1e-4]
```

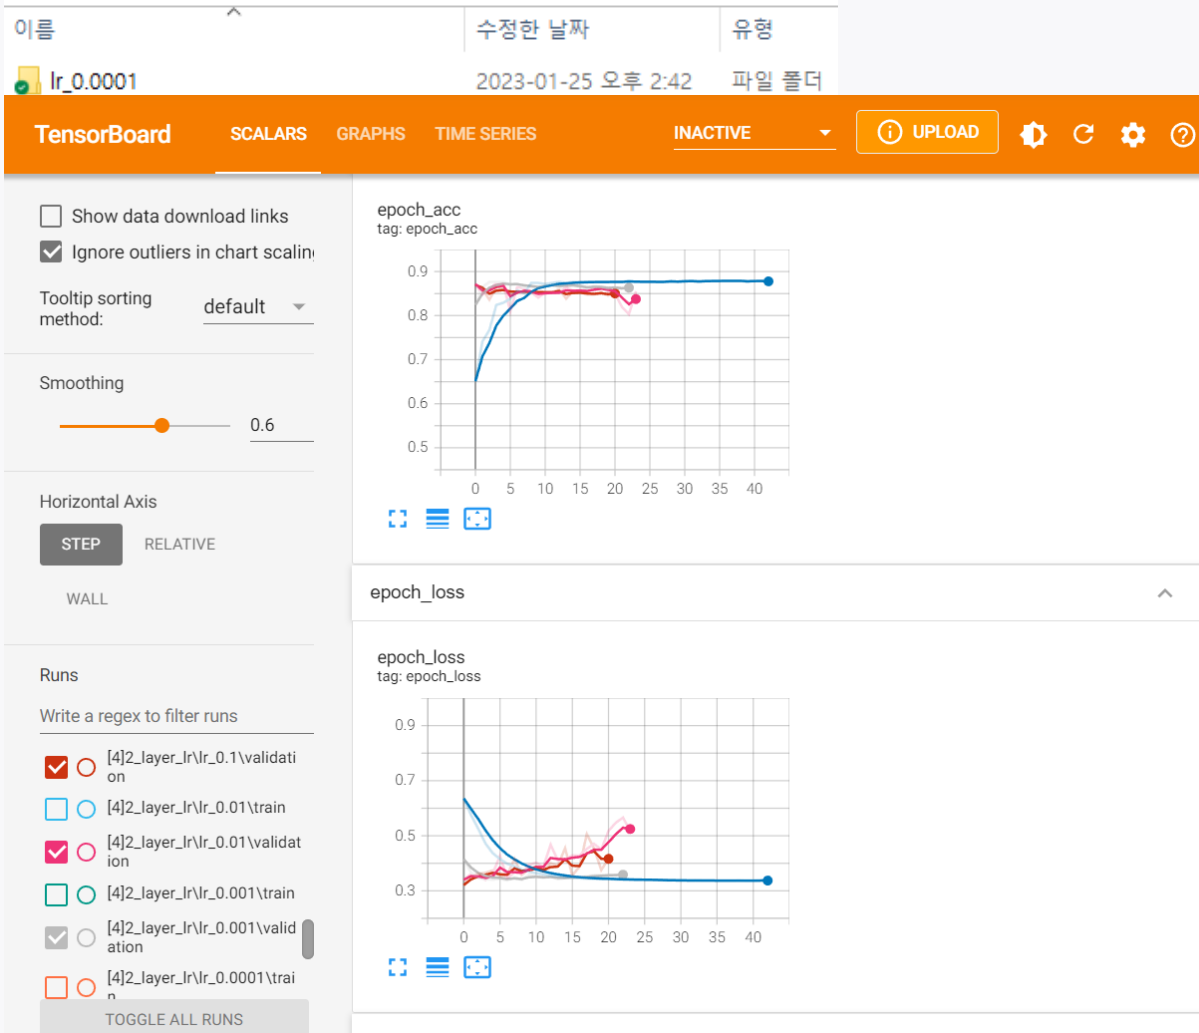
```
for lr in lr_list:
    model = create_model(n_layer=2, lambda_=0)
    opt = Adam(learning_rate=lr)
    model.compile(optimizer=opt,
                  loss='binary_crossentropy',
                  metrics=['acc'])
    logdir = os.path.join('logs', '[4]2_layer_lr', 'lr_{}'.format(lr))
    tb_callback = TensorBoard(logdir)
    es_callback = EarlyStopping(monitor='val_acc', patience=20)
    hist = model.fit(X_train, y_train,
                    validation_split=1/9, epochs = 100,
                    verbose=0, callbacks=[tb_callback, es_callback])
```

```
14/14 [=====] - 0s 614us/step - loss: 0.3471 - acc: 0.8694
14/14 [=====] - 0s 690us/step - loss: 0.4463 - acc: 0.8646
14/14 [=====] - 0s 616us/step - loss: 0.3337 - acc: 0.8575
14/14 [=====] - 0s 537us/step - loss: 0.3281 - acc: 0.8741
```

# Hyperparameter search: Learning rate


```
%tensorboard --logdir=logs/[4]2_layer_lr
```

[현대자동차] > DS30\_5차수 > DAY6 > logs > [4]2\_layer\_lr >



# Hyperparameter search: Learning rate

epoch\_acc  
tag: epoch\_acc



	Name	Smoothed	Value
●	[4]2_layer_lr\lr_0.0001\validation	0.8752	0.8765
●	[4]2_layer_lr\lr_0.001\validation	0.867	0.8646
●	[4]2_layer_lr\lr_0.01\validation	0.8568	0.8575
●	[4]2_layer_lr\lr_0.1\validation	0.8519	0.8527



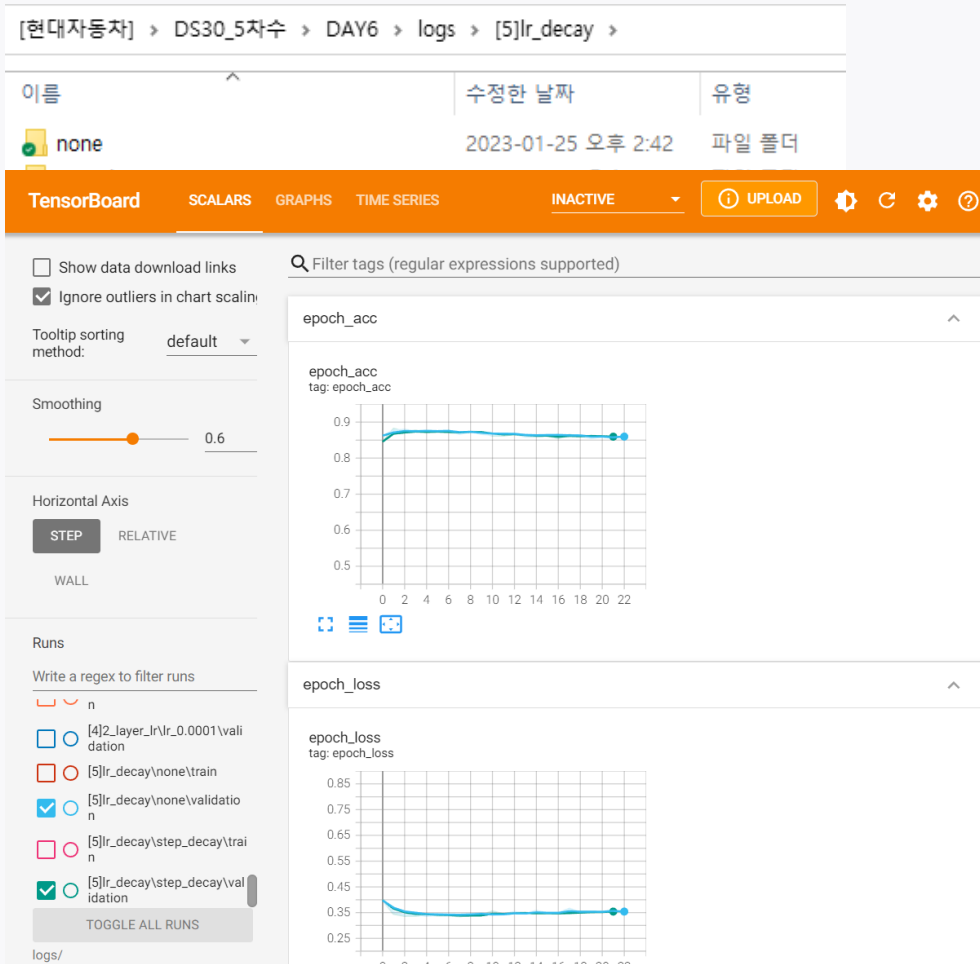
# Effect of learning rate decaying

```
# w/o learning rate decay
model = create_model(n_layer=2, lambda_=0)
opt = Adam(learning_rate=1e-3)
model.compile(optimizer=opt,
              loss='binary_crossentropy',
              metrics=['acc'])
logdir = os.path.join('logs', '[5]lr_decay', 'none')
tb_callback = TensorBoard(logdir)
es_callback = EarlyStopping(monitor='val_acc', patience=20)
hist = model.fit(X_train, y_train,
                validation_split=1/9, epochs = 100,
                verbose=0, callbacks=[tb_callback, es_callback] )
model.evaluate(X_test, y_test)
```

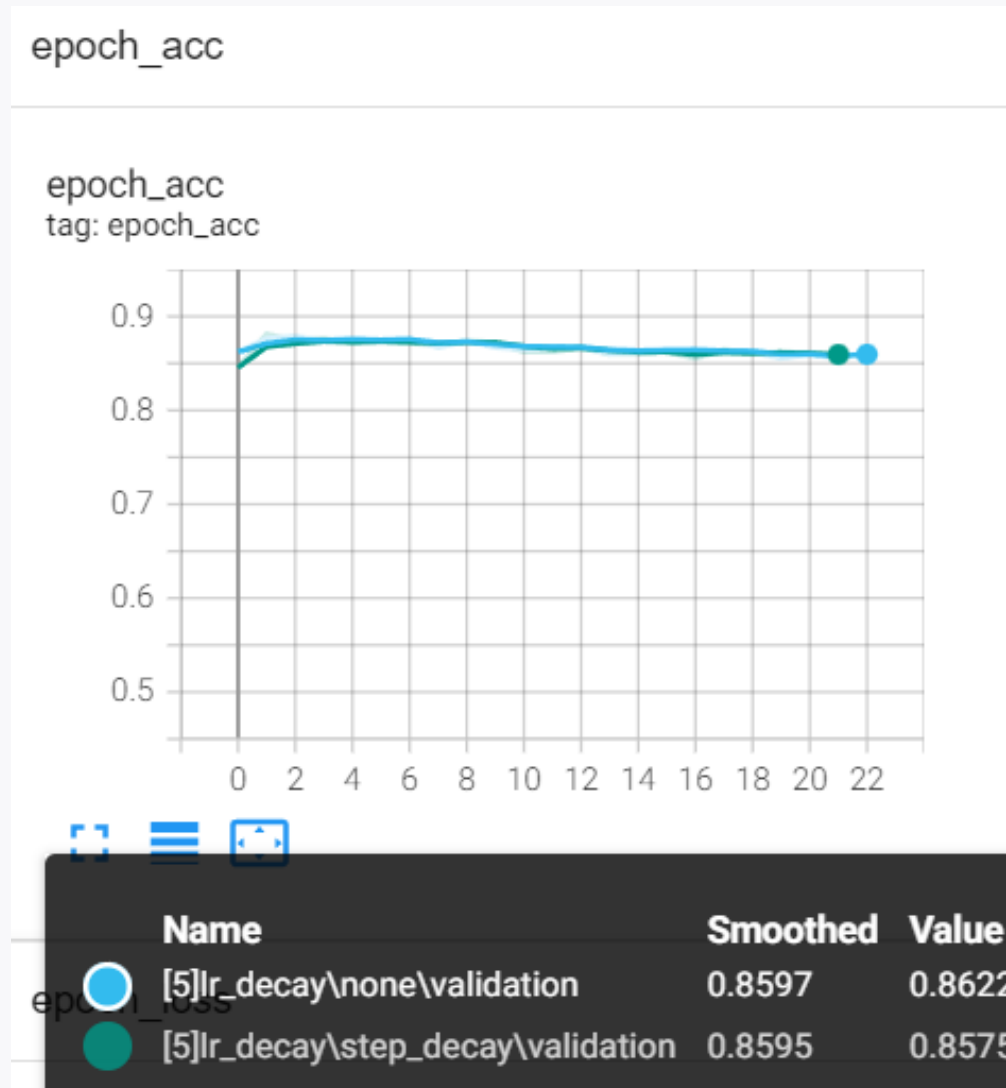
```
# w/ learning rate decay
model = create_model(n_layer=2, lambda_=0)
opt = Adam(learning_rate=1e-3)
model.compile(optimizer=opt,
              loss='binary_crossentropy',
              metrics=['acc'])
logdir = os.path.join('logs', '[5]lr_decay', 'step_decay')
tb_callback = TensorBoard(logdir)
es_callback = EarlyStopping(monitor='val acc', patience=20)
def scheduler(epoch, lr):
    if epoch in [30, 60, 90]:
        lr = lr*0.1
    return lr
lrs_callback = LearningRateScheduler(scheduler)
hist = model.fit(X_train, y_train,
                validation_split=1/9, epochs = 100,
                verbose=0, callbacks=[tb_callback, es_callback, lrs_callback] )
model.evaluate(X_test, y_test)
```

# Effect of learning rate decaying

```
%tensorboard --logdir=logs/[5]lr_decay
```



# Effect of learning rate decaying



# Look ahead

---

Will study how to implement cross validation for a Keras model.