

PS21

January 26, 2023

1 Mini-project #2 (RNN exercise)

1.1 Task: Weather prediction

1.2 Loading Jena climate dataset

```
[1]: import pandas as pd
```

```
[2]: data = pd.read_csv('jena_climate_2009_2016.csv')  
data.describe()
```

```
[2]:
```

	p (mbar)	T (degC)	Tpot (K)	Tdew (degC)	\
count	420551.000000	420551.000000	420551.000000	420551.000000	
mean	989.212776	9.450147	283.492743	4.955854	
std	8.358481	8.423365	8.504471	6.730674	
min	913.600000	-23.010000	250.600000	-25.010000	
25%	984.200000	3.360000	277.430000	0.240000	
50%	989.580000	9.420000	283.470000	5.220000	
75%	994.720000	15.470000	289.530000	10.070000	
max	1015.350000	37.280000	311.340000	23.110000	

	rh (%)	VPmax (mbar)	VPact (mbar)	VPdef (mbar)	\
count	420551.000000	420551.000000	420551.000000	420551.000000	
mean	76.008259	13.576251	9.533756	4.042412	
std	16.476175	7.739020	4.184164	4.896851	
min	12.950000	0.950000	0.790000	0.000000	
25%	65.210000	7.780000	6.210000	0.870000	
50%	79.300000	11.820000	8.860000	2.190000	
75%	89.400000	17.600000	12.350000	5.300000	
max	100.000000	63.770000	28.320000	46.010000	

	sh (g/kg)	H2O _C (mmol/mol)	rho (g/m**3)	wv (m/s)	\
count	420551.000000	420551.000000	420551.000000	420551.000000	
mean	6.022408	9.640223	1216.062748	1.702224	
std	2.656139	4.235395	39.975208	65.446714	
min	0.500000	0.800000	1059.450000	-9999.000000	
25%	3.920000	6.290000	1187.490000	0.990000	
50%	5.590000	8.960000	1213.790000	1.760000	

75%	7.800000	12.490000	1242.770000	2.860000
max	18.130000	28.820000	1393.540000	28.490000

	max. wv (m/s)	wd (deg)
count	420551.000000	420551.000000
mean	3.056555	174.743738
std	69.016932	86.681693
min	-9999.000000	0.000000
25%	1.760000	124.900000
50%	2.960000	198.100000
75%	4.740000	234.100000
max	23.500000	360.000000

1.3 Data preprocessing

Wind speed (and maximum wind speed) is set to -9999.00 for missing entries. Let us fill up the missing entries with the mean.

```
[3]: wv = data['wv (m/s)']
wv_missing_idx = (wv == -9999.00)
wv_mean = wv[~wv_missing_idx].mean()
wv[wv_missing_idx] = wv_mean
```

C:\Users\chsuh\AppData\Local\Temp\ipykernel_23704\832945953.py:4:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
wv[wv_missing_idx] = wv_mean
```

```
[4]: max_wv = data['max. wv (m/s)']
missing_idx = (max_wv == -9999.00)
max_wv_mean = max_wv[~missing_idx].mean()
max_wv[missing_idx] = max_wv_mean
```

C:\Users\chsuh\AppData\Local\Temp\ipykernel_23704\1423667913.py:4:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
max_wv[missing_idx] = max_wv_mean
```

1.4 Remove date_time column

```
[5]: data.pop('Date Time')
data
```

```
[5]:
```

	p (mbar)	T (degC)	Tpot (K)	Tdew (degC)	rh (%)	VPmax (mbar)	\
0	996.52	-8.02	265.40	-8.90	93.30	3.33	
1	996.57	-8.41	265.01	-9.28	93.40	3.23	
2	996.53	-8.51	264.91	-9.31	93.90	3.21	
3	996.51	-8.31	265.12	-9.07	94.20	3.26	
4	996.51	-8.27	265.15	-9.04	94.10	3.27	
...	
420546	1000.07	-4.05	269.10	-8.13	73.10	4.52	
420547	999.93	-3.35	269.81	-8.06	69.71	4.77	
420548	999.82	-3.16	270.01	-8.21	67.91	4.84	
420549	999.81	-4.23	268.94	-8.53	71.80	4.46	
420550	999.82	-4.82	268.36	-8.42	75.70	4.27	
...	
	VPact (mbar)	VPdef (mbar)	sh (g/kg)	H2OC (mmol/mol)	rho (g/m**3)	\	
0	3.11	0.22	1.94		3.12	1307.75	
1	3.02	0.21	1.89		3.03	1309.80	
2	3.01	0.20	1.88		3.02	1310.24	
3	3.07	0.19	1.92		3.08	1309.19	
4	3.08	0.19	1.92		3.09	1309.00	
...		
420546	3.30	1.22	2.06		3.30	1292.98	
420547	3.32	1.44	2.07		3.32	1289.44	
420548	3.28	1.55	2.05		3.28	1288.39	
420549	3.20	1.26	1.99		3.20	1293.56	
420550	3.23	1.04	2.01		3.23	1296.38	
...		
	wv (m/s)	max. wv (m/s)	wd (deg)				
0	1.03	1.75	152.3				
1	0.72	1.50	136.1				
2	0.19	0.63	171.6				
3	0.34	0.50	198.0				
4	0.32	0.63	214.3				
...				
420546	0.67	1.52	240.0				
420547	1.14	1.92	234.3				
420548	1.08	2.00	215.2				
420549	1.49	2.16	225.8				
420550	1.23	1.96	184.9				

```
[420551 rows x 14 columns]
```

1.5 Downsampling

We sample every 60 minutes to construct a dataset for training.

```
[6]: data = data[0::6]
data
```

```
[6]:
```

	p (mbar)	T (degC)	Tpot (K)	Tdew (degC)	rh (%)	VPmax (mbar)	\
0	996.52	-8.02	265.40	-8.90	93.30	3.33	
6	996.50	-7.62	265.81	-8.30	94.80	3.44	
12	996.63	-8.85	264.57	-9.70	93.50	3.12	
18	996.87	-8.84	264.56	-9.69	93.50	3.13	
24	997.05	-9.23	264.15	-10.25	92.20	3.03	
...	
420522	1002.08	-1.40	271.59	-6.10	70.20	5.51	
420528	1001.42	-2.15	270.90	-7.08	68.77	5.21	
420534	1001.05	-2.61	270.47	-6.97	71.80	5.04	
420540	1000.51	-3.22	269.90	-7.63	71.40	4.81	
420546	1000.07	-4.05	269.10	-8.13	73.10	4.52	
	VPact (mbar)	VPdef (mbar)	sh (g/kg)	H2OC (mmol/mol)	rho (g/m**3)	\	
0	3.11	0.22	1.94	3.12	1307.75		
6	3.26	0.18	2.04	3.27	1305.68		
12	2.92	0.20	1.82	2.93	1312.11		
18	2.92	0.20	1.83	2.93	1312.37		
24	2.79	0.24	1.74	2.80	1314.62		
...		
420522	3.87	1.64	2.40	3.86	1282.68		
420528	3.59	1.63	2.23	3.58	1285.50		
420534	3.62	1.42	2.25	3.61	1287.20		
420540	3.44	1.38	2.14	3.44	1289.50		
420546	3.30	1.22	2.06	3.30	1292.98		
	wv (m/s)	max. wv (m/s)	wd (deg)				
0	1.03	1.75	152.3				
6	0.18	0.63	166.5				
12	0.16	0.50	158.3				
18	0.07	0.25	129.3				
24	0.10	0.38	203.9				
...				
420522	1.08	1.68	207.5				
420528	0.79	1.24	184.3				
420534	0.77	1.64	129.1				
420540	0.85	1.54	207.8				
420546	0.67	1.52	240.0				

[70092 rows x 14 columns]

1.6 Take temperature in celsius as label while taking everything as features

```
[7]: features = data
     labels = data[['T (degC)']]
```

```
[8]: print(features.shape)
     print(labels.shape)
```

```
(70092, 14)
```

```
(70092, 1)
```

1.7 Normalization

```
[9]: from sklearn.preprocessing import StandardScaler

     std_scaler = StandardScaler()
     features = std_scaler.fit_transform(features)
```

```
[10]: print(features)
      print(features.shape)
```

```
[[ 0.87420457 -2.07391772 -2.12735513 ... -0.71190538 -0.76237653
   -0.2618485 ]
 [ 0.87181184 -2.02643323 -2.07914744 ... -1.26284569 -1.24217323
   -0.09825609]
 [ 0.8873646  -2.17244806 -2.2249463  ... -1.27580899 -1.29786392
   -0.19272494]
 ...
 [ 1.41615816 -1.43168989 -1.53122591 ... -0.8804283  -0.80949942
   -0.52912624]
 [ 1.35155442 -1.50410375 -1.59824636 ... -0.82857509 -0.85233841
    0.37754438]
 [ 1.29891433 -1.60263408 -1.69231014 ... -0.9452448  -0.86090621
    0.74850745]]
(70092, 14)
```

1.8 Data split (train:val:test=7:2:1)

Let us split chronologically, not randomly.

```
[11]: from sklearn.model_selection import train_test_split

      ## For chronological splitting, we set "shuffle=False"
      X_rest, X_test, y_rest, y_test = train_test_split(features,
                                                         labels,
                                                         test_size=0.1,
                                                         shuffle=False)
      X_train, X_val, y_train, y_val = train_test_split(X_rest,
```

```
y_rest,  
test_size=2/9,  
shuffle=False)
```

```
[12]: print(X_train.shape)  
print(X_val.shape)  
print(X_test.shape)
```

```
(49063, 14)  
(14019, 14)  
(7010, 14)
```

1.9 Set T_{window} and batch size

```
[13]: T = 24  
batch_size = 16
```

1.10 Construct $\{(x_T^{(i)}, y_T^{(i)})\}_{i=1}^{m_T}$

```
[14]: from tensorflow.keras.preprocessing import timeseries_dataset_from_array  
  
# Train batch dataset  
dataset_train = timeseries_dataset_from_array(X_train[:-T],  
                                              y_train[T:],  
                                              sequence_length = T,  
                                              sequence_stride = 1,  
                                              batch_size = batch_size,  
                                              shuffle = True)  
  
# validation batch dataset  
dataset_val = timeseries_dataset_from_array(X_val[:-T],  
                                           y_val[T:],  
                                           sequence_length = T,  
                                           sequence_stride = 1,  
                                           batch_size=batch_size,  
                                           shuffle = False)  
  
# test batch dataset  
dataset_test = timeseries_dataset_from_array(X_test[:-T],  
                                             y_test[T:],  
                                             sequence_length = T,  
                                             sequence_stride = 1,  
                                             batch_size=batch_size,  
                                             shuffle = False)
```

```
[15]: print(type(dataset_train))
```

```
<class 'tensorflow.python.data.ops.dataset_ops.BatchDataset'>
```

```
[16]: print(dataset_train.take)
```

```
<bound method DatasetV2.take of <BatchDataset
element_spec=(TensorSpec(shape=(None, None, 14), dtype=tf.float64, name=None),
TensorSpec(shape=(None, 1), dtype=tf.float64, name=None))>>
```

```
[17]: print(len(dataset_train))
print(len(X_train[:-T])//batch_size)
print(len(dataset_val))
print(len(X_val[:-T])//batch_size)
print(len(dataset_test))
print(len(X_test[:-T])//batch_size)
```

```
3064
3064
874
874
436
436
```

```
[18]: print(dataset_train.take)
print(len(dataset_train))
print(len(dataset_val))
print(len(dataset_test))
```

```
<bound method DatasetV2.take of <BatchDataset
element_spec=(TensorSpec(shape=(None, None, 14), dtype=tf.float64, name=None),
TensorSpec(shape=(None, 1), dtype=tf.float64, name=None))>>
```

```
3064
874
436
```

```
[19]: for batch in dataset_train.take(5):
      inputs, labels = batch
      print(inputs.shape)
      print(labels.shape)
```

```
(16, 24, 14)
(16, 1)
(16, 24, 14)
(16, 1)
(16, 24, 14)
(16, 1)
(16, 24, 14)
(16, 1)
(16, 24, 14)
(16, 1)
```