

Advanced techniques

Lecture 5

Changho Suh

January 23, 2024

Early stopping, dropout, Weight initialization & techniques for training stability

Outline

1. Generalization techniques

Early stopping

Dropout

2. Weight initialization

Xavier's initialization

He's initialization


3. Techniques for training stability

Adam optimizer

Learning rate decaying

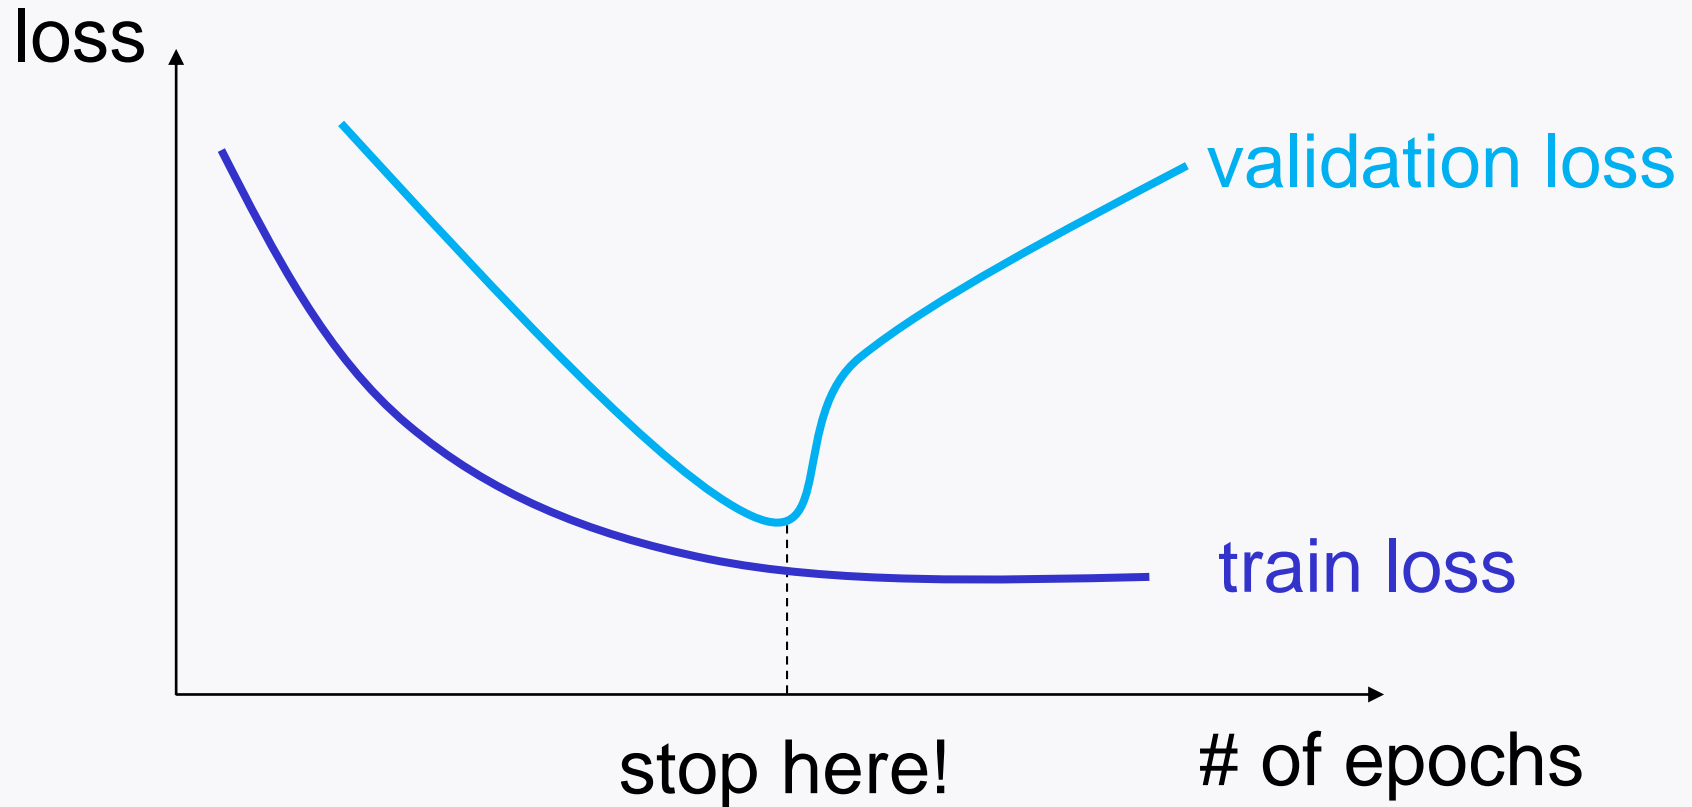
Batch normalization

Early stopping: Motivation



Large # of epochs: **Overfitting** to train data.

Early stopping: Idea



To avoid overfitting: Rely on **validation loss**.


Dropout

$$J(w) = \frac{1}{m_{\mathcal{B}}} \sum_{i=1}^{m_{\mathcal{B}}} \ell(y^{(i)}, \hat{y}^{(i)})$$

In computing a prediction $\hat{y}^{(i)}$ **per example**, construct a random neural network.


How to construct a random neural network?

Construction of a random neural network




Construction of a random neural network

Dropout rate: p (e.g., 0.5)




Construction of a random neural network

Dropout rate: p (e.g., 0.5)



Construction of a random neural network

Dropout rate: p (e.g., 0.5)



Generate this partial NN per example.


Why dropout works?

Indirect effect to incorporate many smaller NNs.


Can interpret the resulting NN as an **averaging ensemble** of all these smaller NNs.

Not overfit to a particular NN; hence generalize better.


Xavier's initialization: Motivation




Xavier's initialization: Motivation



Xavier's initialization: Motivation



Turns out: With $w_{ij}^{[\ell]} \sim \mathcal{N}(0, 1)$,
signals blow up as the network gets deeper.



Xavier's initialization: Motivation

To see this “exploding problem”, consider:

$$z_1^{[\ell]} = \sum_{j=1}^{n^{[\ell-1]}} w_{1j}^{[\ell]} a_j^{[\ell-1]}$$

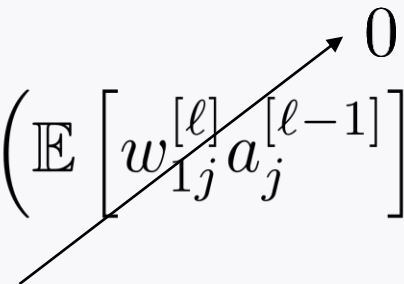
Signal dynamics can be quantified via:

$$\text{var} \left(z_1^{[\ell]} \right) = \text{var} \left(\sum_{j=1}^{n^{[\ell-1]}} w_{1j}^{[\ell]} a_j^{[\ell-1]} \right)$$

Variance computation

$$\text{var} \left(z_1^{[\ell]} \right) = \text{var} \left(\sum_{j=1}^{n^{[\ell-1]}} w_{1j}^{[\ell]} a_j^{[\ell-1]} \right)$$

$$= \sum_{j=1}^{n^{[\ell-1]}} \text{var} \left(w_{1j}^{[\ell]} a_j^{[\ell-1]} \right)$$

$$= \sum_{j=1}^{n^{[\ell-1]}} \mathbb{E} \left[(w_{1j}^{[\ell]})^2 (a_j^{[\ell-1]})^2 \right] - \sum_{j=1}^{n^{[\ell-1]}} \left(\mathbb{E} \left[w_{1j}^{[\ell]} a_j^{[\ell-1]} \right] \right)^2$$


Assumption:

- (i) weights independent
- (ii) input independent
- (iii) weights/input ind.
- (iv) zero mean

Variance computation

$$\begin{aligned}
 \text{var} \left(z_1^{[\ell]} \right) &= \sum_{j=1}^{n^{[\ell-1]}} \mathbb{E} \left[(w_{1j}^{[\ell]})^2 (a_j^{[\ell-1]})^2 \right] \\
 &= \sum_{j=1}^{n^{[\ell-1]}} \mathbb{E} \left[(w_{1j}^{[\ell]})^2 \right] \mathbb{E} \left[(a_j^{[\ell-1]})^2 \right] \\
 &= \sum_{j=1}^{n^{[\ell-1]}} \text{var} \left(w_{1j}^{[\ell]} \right) \text{var} \left(a_j^{[\ell-1]} \right)
 \end{aligned}$$

Assumption:

- (i) weights independent
- (ii) input independent
- (iii) weights/input ind.
- (iv) zero mean

Exploding problem

$$\text{var} \left(z_1^{[\ell]} \right) = \sum_{j=1}^{n^{[\ell-1]}} \text{var} \left(w_{1j}^{[\ell]} \right) \text{var} \left(a_j^{[\ell-1]} \right)$$

Suppose: $\text{var} \left(a_j^{[\ell-1]} \right) = 1, \text{var} \left(w_{1j}^{[\ell]} \right) = 1$

Then: $\text{var} \left(z_1^{[\ell]} \right) = n^{[\ell-1]}$

As the network gets deeper, explode!

Xavier's initialization

$$\text{var} \left(z_1^{[\ell]} \right) = \sum_{j=1}^{n^{[\ell-1]}} \text{var} \left(w_{1j}^{[\ell]} \right) \text{var} \left(a_j^{[\ell-1]} \right)$$

Suppose: $\text{var} \left(a_j^{[\ell-1]} \right) = 1$

Idea: Set $\text{var} \left(w_{1j}^{[\ell]} \right) = \frac{1}{n^{[\ell-1]}}$

$$w_{ij}^{[\ell]} \text{ i.i.d. } \sim \mathcal{N} \left(0, \frac{1}{n^{[\ell-1]}} \right)$$

He's initialization: Motivation

$$\begin{aligned} a_1^{[\ell]} &= \text{ReLU}(z_1^{[\ell]}) \\ &= \max(0, z_1^{[\ell]}) \end{aligned}$$

$$\text{var} \left(a_1^{[\ell]} \right) = \frac{1}{2} \text{var} \left(z_1^{[\ell]} \right)$$

Xavier's initialization $w_{ij}^{[\ell]} \sim \mathcal{N} \left(0, \frac{1}{n^{[\ell-1]}} \right)$

$$\longrightarrow \text{var} \left(a_1^{[\ell]} \right) = \frac{1}{2} \text{var} \left(a_j^{[\ell-1]} \right)$$

He's initialization

$$w_{ij}^{[\ell]} \sim \mathcal{N} \left(0, \frac{2}{n^{[\ell-1]}} \right)$$

Techniques for training stability:

Adam optimizer

Learning rate decaying


Batch normalization

Learning rate decaying

Three popular choices:

1. $\alpha^{(t)} = \gamma^t \quad 0 < \gamma < 1$

2. $\alpha^{(t)} = \frac{1}{\sqrt{t}}$



Batch normalization: Motivation

Turns out: Different signal scalings across distinct layers incur training instability.

One prominent way to address this:

Batch normalization


Batch

Recall the cost function used for gradient descent:

$$J(w^{(t)}) := \frac{1}{m} \sum_{i=1}^m -y^{(i)} \log \hat{y}^{(i)} - (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})$$

Issue: Computationally heavy for a **large m** .


Hence: In practice, use a chunk of examples, called a *batch*.



Batch normalization

A hidden layer

BN



1. Normalization

$$z_{\text{norm}} = \frac{z - \mu_{\mathcal{B}}}{\sigma_{\mathcal{B}}}$$

$$\mu_{\mathcal{B}} = \frac{1}{m_{\mathcal{B}}} \sum_{i \in \mathcal{B}} z^{(i)} \quad \sigma_{\mathcal{B}}^2 = \frac{1}{m_{\mathcal{B}}} \sum_{i \in \mathcal{B}} (z^{(i)} - \mu_{\mathcal{B}})^2$$

2. Customized scaling

$$\tilde{z} = \underset{\substack{\uparrow \\ \text{learnable parameters}}}{\gamma} z_{\text{norm}} + \beta$$

learnable parameters



Look ahead

Will study:

hyperparameter search

cross validation