

PS5

January 15, 2024

```
[1]: from shap.datasets import adult
```

```
X, y = adult()
```

```
print(X)
```

```
print(y)
```

	Age	Workclass	Education-Num	Marital Status	Occupation \
0	39.0	7	13.0	4	1
1	50.0	6	13.0	2	4
2	38.0	4	9.0	0	6
3	53.0	4	7.0	2	6
4	28.0	4	13.0	2	10
...
32556	27.0	4	12.0	2	13
32557	40.0	4	9.0	2	7
32558	58.0	4	9.0	6	1
32559	22.0	4	9.0	4	1
32560	52.0	5	9.0	2	4

	Relationship	Race	Sex	Capital Gain	Capital Loss	Hours per week \
0	0	4	1	2174.0	0.0	40.0
1	4	4	1	0.0	0.0	13.0
2	0	4	1	0.0	0.0	40.0
3	4	2	1	0.0	0.0	40.0
4	5	2	0	0.0	0.0	40.0
...
32556	5	4	0	0.0	0.0	38.0
32557	4	4	1	0.0	0.0	40.0
32558	1	4	0	0.0	0.0	40.0
32559	3	4	1	0.0	0.0	20.0
32560	5	4	0	15024.0	0.0	40.0

	Country
0	39
1	39
2	39
3	39

```

4          5
...      ...
32556     39
32557     39
32558     39
32559     39
32560     39

```

```

[32561 rows x 12 columns]
[False False False ... False False  True]

```

```

[2]: numerical_columns = ['Age', 'Education-Num', 'Capital Gain', 'Capital Loss', 'Hours_
    ↪per week']
    categorical_columns = ['Workclass', 'Marital_
    ↪Status', 'Occupation', 'Relationship', 'Race', 'Sex', 'Country']

```

0.1 Conversion of categorical data

```

[3]: import pandas as pd # for one-hot encoding

    from sklearn.preprocessing import StandardScaler # for normalization

```

```

[4]: # Normalization of numerical data
    for column in numerical_columns:
        scaler = StandardScaler()
        X[column] = scaler.fit_transform(X[column].values.reshape(-1,1))

    print(X)

```

	Age	Workclass	Education-Num	Marital Status	Occupation \
0	0.030671	7	1.134739	4	1
1	0.837109	6	1.134739	2	4
2	-0.042642	4	-0.420060	0	6
3	1.057047	4	-1.197459	2	6
4	-0.775768	4	1.134739	2	10
...
32556	-0.849080	4	0.746039	2	13
32557	0.103983	4	-0.420060	2	7
32558	1.423610	4	-0.420060	6	1
32559	-1.215643	4	-0.420060	4	1
32560	0.983734	5	-0.420060	2	4

	Relationship	Race	Sex	Capital Gain	Capital Loss	Hours per week \
0	0	4	1	0.148453	-0.21666	-0.035429
1	4	4	1	-0.145920	-0.21666	-2.222153
2	0	4	1	-0.145920	-0.21666	-0.035429
3	4	2	1	-0.145920	-0.21666	-0.035429
4	5	2	0	-0.145920	-0.21666	-0.035429

...	
32556		5	4	0	-0.145920	-0.21666	-0.197409
32557		4	4	1	-0.145920	-0.21666	-0.035429
32558		1	4	0	-0.145920	-0.21666	-0.035429
32559		3	4	1	-0.145920	-0.21666	-1.655225
32560		5	4	0	1.888424	-0.21666	-0.035429

	Country
0	39
1	39
2	39
3	39
4	5

...	...
32556	39
32557	39
32558	39
32559	39
32560	39

[32561 rows x 12 columns]

```
[5]: # Data type change of categorical data
for column in categorical_columns:
    X[column] = X[column].astype('category')

print(X)
```

	Age	Workclass	Education-Num	Marital Status	Occupation \
0	0.030671	7	1.134739	4	1
1	0.837109	6	1.134739	2	4
2	-0.042642	4	-0.420060	0	6
3	1.057047	4	-1.197459	2	6
4	-0.775768	4	1.134739	2	10

...
32556	-0.849080	4	0.746039	2	13
32557	0.103983	4	-0.420060	2	7
32558	1.423610	4	-0.420060	6	1
32559	-1.215643	4	-0.420060	4	1
32560	0.983734	5	-0.420060	2	4

	Relationship	Race	Sex	Capital Gain	Capital Loss	Hours per week \	
0		0	4	1	0.148453	-0.21666	-0.035429
1		4	4	1	-0.145920	-0.21666	-2.222153
2		0	4	1	-0.145920	-0.21666	-0.035429
3		4	2	1	-0.145920	-0.21666	-0.035429
4		5	2	0	-0.145920	-0.21666	-0.035429

...
-----	-----	-----	----	-----	-----	-----

32556	5	4	0	-0.145920	-0.21666	-0.197409
32557	4	4	1	-0.145920	-0.21666	-0.035429
32558	1	4	0	-0.145920	-0.21666	-0.035429
32559	3	4	1	-0.145920	-0.21666	-1.655225
32560	5	4	0	1.888424	-0.21666	-0.035429

	Country
0	39
1	39
2	39
3	39
4	5
...	...
32556	39
32557	39
32558	39
32559	39
32560	39

[32561 rows x 12 columns]

```
[6]: # One-hot encoding of categorical data
X = pd.get_dummies(X)

# Conversion of data frame to numpy
X = X.values

# Converision: {False, True} --> {0., 1.}
y = y.astype(float)
```

```
[7]: print(X.shape)
print(y.shape)
print(y)
```

```
(32561, 91)
(32561,)
[0. 0. 0. ... 0. 0. 1.]
```

0.2 train-val-test split

```
[8]: from sklearn.model_selection import train_test_split

X_,X_test,y_,y_test = train_test_split(X,y,test_size=1/10,stratify=y)
X_train,X_val,y_train,y_val = train_test_split(X_,y_,test_size=1/9,stratify=y_)

print(X_train.shape)
print(X_val.shape)
print(X_test.shape)
```

```
(26048, 91)
(3256, 91)
(3257, 91)
```

0.3 Logistic regression

```
[9]: from sklearn.linear_model import LogisticRegression
```

```
[10]: model_LR = LogisticRegression()

# training
model_LR.fit(X_train, y_train)

# evaluation
val_acc = model_LR.score(X_val, y_val)

print(val_acc)
```

```
0.8541154791154791
```

```
lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
[11]: from joblib import dump

dump(model_LR, 'LR_sample.joblib')
```

```
[11]: ['LR_sample.joblib']
```

```
[12]: from joblib import load

load('LR_sample.joblib')
```

```
[12]: LogisticRegression()
```

0.4 A 2-layer DNN

```
[13]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import Adam

model = Sequential()
model.add(Dense(128, activation='relu'))
```

```

model.add(Dense(1, activation='sigmoid'))

opt = Adam(learning_rate=0.01,
           beta_1 = 0.9,
           beta_2 = 0.999)

#model.compile(optimizer=opt,
#              loss='binary_crossentropy')

model.compile(optimizer=opt,
              loss='binary_crossentropy',
              metrics=['acc'])

model.fit(X_train,y_train, epochs=10)

```

```

Epoch 1/10
814/814 [=====] - 3s 2ms/step - loss: 0.3232 - acc:
0.8516
Epoch 2/10
814/814 [=====] - 2s 2ms/step - loss: 0.3138 - acc:
0.8529
Epoch 3/10
814/814 [=====] - 2s 2ms/step - loss: 0.3065 - acc:
0.8579
Epoch 4/10
814/814 [=====] - 2s 3ms/step - loss: 0.3027 - acc:
0.8590
Epoch 5/10
814/814 [=====] - 2s 3ms/step - loss: 0.2979 - acc:
0.8612
Epoch 6/10
814/814 [=====] - 2s 3ms/step - loss: 0.2940 - acc:
0.8613
Epoch 7/10
814/814 [=====] - 2s 3ms/step - loss: 0.2905 - acc:
0.8660
Epoch 8/10
814/814 [=====] - 2s 2ms/step - loss: 0.2864 - acc:
0.8664
Epoch 9/10
814/814 [=====] - 2s 3ms/step - loss: 0.2831 - acc:
0.8674
Epoch 10/10
814/814 [=====] - 2s 3ms/step - loss: 0.2817 - acc:
0.8689

```

```
[13]: <keras.callbacks.History at 0x1a10329ba30>
```

```
[14]: val_hist = model.evaluate(X_val,y_val)

print(val_hist)
```

```
102/102 [=====] - 1s 4ms/step - loss: 0.3328 - acc:
0.8600
[0.3327931761741638, 0.859950840473175]
```

0.5 Regularization

```
[15]: from tensorflow.keras.regularizers import l2

model = Sequential()

model.add(Dense(128,
                kernel_regularizer=l2(0.01),
                bias_regularizer=l2(0.01)))
```

0.6 Early stopping

```
[16]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping

model = Sequential()
model.add(Dense(128,kernel_regularizer=l2(0.01),
                bias_regularizer=l2(0.01),
                activation='relu'))
model.add(Dense(1,kernel_regularizer=l2(0.01),
                bias_regularizer=l2(0.01),
                activation='sigmoid'))
opt = Adam(learning_rate=0.01,beta_1 = 0.9,beta_2 = 0.999)
model.compile(optimizer=opt,
              loss='binary_crossentropy')

#model.compile(optimizer=opt,
#              loss='binary_crossentropy',
#              metrics=['acc'])

#es_callback = EarlyStopping(monitor='val_acc', patience=15)
#es_callback = EarlyStopping(monitor='val_loss', patience=15)
es_callback = EarlyStopping(monitor='val_loss', patience=15)

hist = model.fit(X_train, y_train,
```

```
validation_data=(X_val, y_val),  
epochs=100, callbacks=[es_callback])
```

```
Epoch 1/100  
814/814 [=====] - 4s 4ms/step - loss: 0.4216 -  
val_loss: 0.3996  
Epoch 2/100  
814/814 [=====] - 3s 3ms/step - loss: 0.4030 -  
val_loss: 0.4026  
Epoch 3/100  
814/814 [=====] - 2s 2ms/step - loss: 0.4030 -  
val_loss: 0.3996  
Epoch 4/100  
814/814 [=====] - 2s 2ms/step - loss: 0.3999 -  
val_loss: 0.4011  
Epoch 5/100  
814/814 [=====] - 2s 2ms/step - loss: 0.4002 -  
val_loss: 0.4001  
Epoch 6/100  
814/814 [=====] - 2s 3ms/step - loss: 0.3978 -  
val_loss: 0.4024  
Epoch 7/100  
814/814 [=====] - 2s 2ms/step - loss: 0.3986 -  
val_loss: 0.3934  
Epoch 8/100  
814/814 [=====] - 2s 3ms/step - loss: 0.3986 -  
val_loss: 0.3961  
Epoch 9/100  
814/814 [=====] - 2s 2ms/step - loss: 0.3975 -  
val_loss: 0.4015  
Epoch 10/100  
814/814 [=====] - 2s 3ms/step - loss: 0.3978 -  
val_loss: 0.3999  
Epoch 11/100  
814/814 [=====] - 2s 3ms/step - loss: 0.3978 -  
val_loss: 0.3975  
Epoch 12/100  
814/814 [=====] - 2s 3ms/step - loss: 0.3985 -  
val_loss: 0.3974  
Epoch 13/100  
814/814 [=====] - 2s 3ms/step - loss: 0.3966 -  
val_loss: 0.3942  
Epoch 14/100  
814/814 [=====] - 2s 2ms/step - loss: 0.3981 -  
val_loss: 0.3993  
Epoch 15/100  
814/814 [=====] - 2s 2ms/step - loss: 0.3978 -  
val_loss: 0.3936
```



```

Epoch 16/100
814/814 [=====] - 2s 2ms/step - loss: 0.3972 -
val_loss: 0.3985
Epoch 17/100
814/814 [=====] - 2s 3ms/step - loss: 0.3979 -
val_loss: 0.3960
Epoch 18/100
814/814 [=====] - 3s 3ms/step - loss: 0.4002 -
val_loss: 0.3960
Epoch 19/100
814/814 [=====] - 2s 2ms/step - loss: 0.3981 -
val_loss: 0.4107
Epoch 20/100
814/814 [=====] - 2s 3ms/step - loss: 0.4000 -
val_loss: 0.4014
Epoch 21/100
814/814 [=====] - 2s 3ms/step - loss: 0.3972 -
val_loss: 0.3948
Epoch 22/100
814/814 [=====] - 2s 2ms/step - loss: 0.3963 -
val_loss: 0.3941

```


```

[17]: train_loss = hist.history['loss']
      val_loss = hist.history['val_loss']
      #train_acc = hist.history['acc']
      #val_acc = hist.history['val_acc']

      import matplotlib.pyplot as plt

      plt.plot(train_loss,label='train loss')
      plt.plot(val_loss,label='val loss')
      plt.legend()
      plt.show()

```



0.7 Dropout

```
[18]: from tensorflow.keras.layers import Dropout

model = Sequential()
model.add(Dense(128, kernel_regularizer=l2(0.01),
                bias_regularizer=l2(0.01),
                activation='relu'))

model.add(Dropout(0.9))
```

0.8 Weight initialization

```
[19]: from tensorflow.keras.initializers import HeNormal

init = HeNormal()

model.add(Dense(128, kernel_regularizer=l2(0.01),
                bias_regularizer=l2(0.01),
                kernel_initializer=init,
                activation='relu'))
```

0.9 Batch normalization

```
[20]: from tensorflow.keras.layers import BatchNormalization
      from tensorflow.keras.layers import ReLU

      init = HeNormal()

      model.add(Dense(128, kernel_regularizer=l2(0.01),
                      bias_regularizer=l2(0.01),
                      kernel_initializer=init))

      model.add(BatchNormalization())

      model.add(ReLU())
```

0.10 Learning rate decaying

```
[21]: from tensorflow.keras.initializers import HeNormal
      from tensorflow.keras.layers import BatchNormalization
      from tensorflow.keras.layers import ReLU
      from tensorflow.keras.layers import Dropout
      from tensorflow.keras.callbacks import EarlyStopping
      from tensorflow.keras.callbacks import LearningRateScheduler

      init = HeNormal()

      model = Sequential()
      model.add(Dense(128, kernel_regularizer=l2(0.01),
                      bias_regularizer=l2(0.01),
                      kernel_initializer=init))
      model.add(BatchNormalization())
      model.add(ReLU())
      model.add(Dropout(0.5))
      model.add(Dense(1, activation='sigmoid'))

      opt = Adam(learning_rate=0.01, beta_1 = 0.9, beta_2 = 0.999)
      #model.compile(optimizer=opt,
      #               loss='binary_crossentropy',
      #               metrics=['acc'])
      model.compile(optimizer=opt,
                   loss='binary_crossentropy')

      es_callback = EarlyStopping(monitor='val_loss', patience=15)
      #es_callback = EarlyStopping(monitor='val_acc', patience=15)

      def scheduler(epoch, lr):
          if epoch in [20, 40, 60]:
```

```

        lr = 0.1*lr
    else:
        lr = lr
    return lr

ls_callback = LearningRateScheduler(scheduler)

hist = model.fit(X_train, y_train,
                 validation_data=(X_val, y_val),
                 epochs=100, callbacks=[es_callback,ls_callback])

```

```

Epoch 1/100
814/814 [=====] - 3s 3ms/step - loss: 0.4951 -
val_loss: 0.4659 - lr: 0.0100
Epoch 2/100
814/814 [=====] - 2s 3ms/step - loss: 0.4420 -
val_loss: 0.4038 - lr: 0.0100
Epoch 3/100
814/814 [=====] - 2s 3ms/step - loss: 0.4307 -
val_loss: 0.4165 - lr: 0.0100
Epoch 4/100
814/814 [=====] - 3s 3ms/step - loss: 0.4267 -
val_loss: 0.4006 - lr: 0.0100
Epoch 5/100
814/814 [=====] - 2s 3ms/step - loss: 0.4259 -
val_loss: 0.4203 - lr: 0.0100
Epoch 6/100
814/814 [=====] - 2s 3ms/step - loss: 0.4253 -
val_loss: 0.4111 - lr: 0.0100
Epoch 7/100
814/814 [=====] - 2s 3ms/step - loss: 0.4245 -
val_loss: 0.3966 - lr: 0.0100
Epoch 8/100
814/814 [=====] - 3s 3ms/step - loss: 0.4228 -
val_loss: 0.3943 - lr: 0.0100
Epoch 9/100
814/814 [=====] - 3s 3ms/step - loss: 0.4227 -
val_loss: 0.3920 - lr: 0.0100
Epoch 10/100
814/814 [=====] - 3s 3ms/step - loss: 0.4204 -
val_loss: 0.4045 - lr: 0.0100
Epoch 11/100
814/814 [=====] - 3s 4ms/step - loss: 0.4207 -
val_loss: 0.4148 - lr: 0.0100
Epoch 12/100
814/814 [=====] - 3s 3ms/step - loss: 0.4207 -
val_loss: 0.3959 - lr: 0.0100

```

Epoch 13/100
814/814 [=====] - 3s 3ms/step - loss: 0.4214 -
val_loss: 0.4031 - lr: 0.0100
Epoch 14/100
814/814 [=====] - 3s 3ms/step - loss: 0.4186 -
val_loss: 0.4039 - lr: 0.0100
Epoch 15/100
814/814 [=====] - 2s 3ms/step - loss: 0.4201 -
val_loss: 0.3973 - lr: 0.0100
Epoch 16/100
814/814 [=====] - 3s 4ms/step - loss: 0.4164 -
val_loss: 0.3890 - lr: 0.0100
Epoch 17/100
814/814 [=====] - 3s 3ms/step - loss: 0.4185 -
val_loss: 0.4017 - lr: 0.0100
Epoch 18/100
814/814 [=====] - 2s 3ms/step - loss: 0.4198 -
val_loss: 0.3992 - lr: 0.0100
Epoch 19/100
814/814 [=====] - 2s 3ms/step - loss: 0.4193 -
val_loss: 0.3931 - lr: 0.0100
Epoch 20/100
814/814 [=====] - 3s 3ms/step - loss: 0.4203 -
val_loss: 0.3863 - lr: 0.0100
Epoch 21/100
814/814 [=====] - 3s 4ms/step - loss: 0.3694 -
val_loss: 0.3417 - lr: 1.0000e-03
Epoch 22/100
814/814 [=====] - 3s 4ms/step - loss: 0.3557 -
val_loss: 0.3389 - lr: 1.0000e-03
Epoch 23/100
814/814 [=====] - 3s 4ms/step - loss: 0.3487 -
val_loss: 0.3373 - lr: 1.0000e-03
Epoch 24/100
814/814 [=====] - 3s 3ms/step - loss: 0.3488 -
val_loss: 0.3351 - lr: 1.0000e-03
Epoch 25/100
814/814 [=====] - 3s 4ms/step - loss: 0.3455 -
val_loss: 0.3318 - lr: 1.0000e-03
Epoch 26/100
814/814 [=====] - 3s 4ms/step - loss: 0.3476 -
val_loss: 0.3328 - lr: 1.0000e-03
Epoch 27/100
814/814 [=====] - 3s 4ms/step - loss: 0.3471 -
val_loss: 0.3327 - lr: 1.0000e-03
Epoch 28/100
814/814 [=====] - 3s 3ms/step - loss: 0.3477 -
val_loss: 0.3393 - lr: 1.0000e-03

Epoch 29/100
814/814 [=====] - 3s 4ms/step - loss: 0.3475 -
val_loss: 0.3408 - lr: 1.0000e-03
Epoch 30/100
814/814 [=====] - 4s 4ms/step - loss: 0.3467 -
val_loss: 0.3379 - lr: 1.0000e-03
Epoch 31/100
814/814 [=====] - 3s 3ms/step - loss: 0.3486 -
val_loss: 0.3344 - lr: 1.0000e-03
Epoch 32/100
814/814 [=====] - 3s 4ms/step - loss: 0.3470 -
val_loss: 0.3319 - lr: 1.0000e-03
Epoch 33/100
814/814 [=====] - 2s 3ms/step - loss: 0.3447 -
val_loss: 0.3334 - lr: 1.0000e-03
Epoch 34/100
814/814 [=====] - 3s 3ms/step - loss: 0.3474 -
val_loss: 0.3366 - lr: 1.0000e-03
Epoch 35/100
814/814 [=====] - 3s 3ms/step - loss: 0.3471 -
val_loss: 0.3358 - lr: 1.0000e-03
Epoch 36/100
814/814 [=====] - 2s 3ms/step - loss: 0.3481 -
val_loss: 0.3347 - lr: 1.0000e-03
Epoch 37/100
814/814 [=====] - 2s 3ms/step - loss: 0.3476 -
val_loss: 0.3300 - lr: 1.0000e-03
Epoch 38/100
814/814 [=====] - 3s 3ms/step - loss: 0.3475 -
val_loss: 0.3356 - lr: 1.0000e-03
Epoch 39/100
814/814 [=====] - 2s 2ms/step - loss: 0.3478 -
val_loss: 0.3344 - lr: 1.0000e-03
Epoch 40/100
814/814 [=====] - 3s 3ms/step - loss: 0.3463 -
val_loss: 0.3370 - lr: 1.0000e-03
Epoch 41/100
814/814 [=====] - 2s 3ms/step - loss: 0.3403 -
val_loss: 0.3296 - lr: 1.0000e-04
Epoch 42/100
814/814 [=====] - 3s 4ms/step - loss: 0.3388 -
val_loss: 0.3290 - lr: 1.0000e-04
Epoch 43/100
814/814 [=====] - 3s 4ms/step - loss: 0.3372 -
val_loss: 0.3275 - lr: 1.0000e-04
Epoch 44/100
814/814 [=====] - 3s 4ms/step - loss: 0.3345 -
val_loss: 0.3266 - lr: 1.0000e-04

Epoch 45/100
814/814 [=====] - 3s 4ms/step - loss: 0.3292 -
val_loss: 0.3270 - lr: 1.0000e-04
Epoch 46/100
814/814 [=====] - 3s 4ms/step - loss: 0.3307 -
val_loss: 0.3248 - lr: 1.0000e-04
Epoch 47/100
814/814 [=====] - 3s 4ms/step - loss: 0.3317 -
val_loss: 0.3250 - lr: 1.0000e-04
Epoch 48/100
814/814 [=====] - 3s 3ms/step - loss: 0.3315 -
val_loss: 0.3237 - lr: 1.0000e-04
Epoch 49/100
814/814 [=====] - 3s 3ms/step - loss: 0.3316 -
val_loss: 0.3243 - lr: 1.0000e-04
Epoch 50/100
814/814 [=====] - 3s 4ms/step - loss: 0.3311 -
val_loss: 0.3255 - lr: 1.0000e-04
Epoch 51/100
814/814 [=====] - 3s 3ms/step - loss: 0.3298 -
val_loss: 0.3232 - lr: 1.0000e-04
Epoch 52/100
814/814 [=====] - 3s 4ms/step - loss: 0.3276 -
val_loss: 0.3222 - lr: 1.0000e-04
Epoch 53/100
814/814 [=====] - 3s 3ms/step - loss: 0.3289 -
val_loss: 0.3236 - lr: 1.0000e-04
Epoch 54/100
814/814 [=====] - 3s 4ms/step - loss: 0.3294 -
val_loss: 0.3214 - lr: 1.0000e-04
Epoch 55/100
814/814 [=====] - 3s 4ms/step - loss: 0.3288 -
val_loss: 0.3223 - lr: 1.0000e-04
Epoch 56/100
814/814 [=====] - 3s 3ms/step - loss: 0.3297 -
val_loss: 0.3215 - lr: 1.0000e-04
Epoch 57/100
814/814 [=====] - 3s 4ms/step - loss: 0.3294 -
val_loss: 0.3231 - lr: 1.0000e-04
Epoch 58/100
814/814 [=====] - 3s 3ms/step - loss: 0.3271 -
val_loss: 0.3204 - lr: 1.0000e-04
Epoch 59/100
814/814 [=====] - 3s 4ms/step - loss: 0.3256 -
val_loss: 0.3211 - lr: 1.0000e-04
Epoch 60/100
814/814 [=====] - 2s 3ms/step - loss: 0.3267 -
val_loss: 0.3205 - lr: 1.0000e-04


Epoch 61/100
814/814 [=====] - 3s 3ms/step - loss: 0.3253 -
val_loss: 0.3200 - lr: 1.0000e-05
Epoch 62/100
814/814 [=====] - 3s 4ms/step - loss: 0.3269 -
val_loss: 0.3222 - lr: 1.0000e-05
Epoch 63/100
814/814 [=====] - 3s 4ms/step - loss: 0.3228 -
val_loss: 0.3213 - lr: 1.0000e-05
Epoch 64/100
814/814 [=====] - 3s 3ms/step - loss: 0.3237 -
val_loss: 0.3213 - lr: 1.0000e-05
Epoch 65/100
814/814 [=====] - 3s 3ms/step - loss: 0.3260 -
val_loss: 0.3216 - lr: 1.0000e-05
Epoch 66/100
814/814 [=====] - 3s 4ms/step - loss: 0.3249 -
val_loss: 0.3204 - lr: 1.0000e-05
Epoch 67/100
814/814 [=====] - 3s 4ms/step - loss: 0.3249 -
val_loss: 0.3209 - lr: 1.0000e-05
Epoch 68/100
814/814 [=====] - 3s 4ms/step - loss: 0.3240 -
val_loss: 0.3198 - lr: 1.0000e-05
Epoch 69/100
814/814 [=====] - 3s 4ms/step - loss: 0.3260 -
val_loss: 0.3218 - lr: 1.0000e-05
Epoch 70/100
814/814 [=====] - 3s 3ms/step - loss: 0.3249 -
val_loss: 0.3210 - lr: 1.0000e-05
Epoch 71/100
814/814 [=====] - 3s 3ms/step - loss: 0.3267 -
val_loss: 0.3205 - lr: 1.0000e-05
Epoch 72/100
814/814 [=====] - 3s 4ms/step - loss: 0.3216 -
val_loss: 0.3220 - lr: 1.0000e-05
Epoch 73/100
814/814 [=====] - 2s 3ms/step - loss: 0.3207 -
val_loss: 0.3200 - lr: 1.0000e-05
Epoch 74/100
814/814 [=====] - 3s 4ms/step - loss: 0.3244 -
val_loss: 0.3211 - lr: 1.0000e-05
Epoch 75/100
814/814 [=====] - 3s 3ms/step - loss: 0.3245 -
val_loss: 0.3216 - lr: 1.0000e-05
Epoch 76/100
814/814 [=====] - 3s 3ms/step - loss: 0.3236 -
val_loss: 0.3193 - lr: 1.0000e-05

Epoch 77/100
814/814 [=====] - 3s 4ms/step - loss: 0.3239 -
val_loss: 0.3218 - lr: 1.0000e-05
Epoch 78/100
814/814 [=====] - 3s 4ms/step - loss: 0.3217 -
val_loss: 0.3202 - lr: 1.0000e-05
Epoch 79/100
814/814 [=====] - 3s 3ms/step - loss: 0.3226 -
val_loss: 0.3205 - lr: 1.0000e-05
Epoch 80/100
814/814 [=====] - 3s 4ms/step - loss: 0.3245 -
val_loss: 0.3216 - lr: 1.0000e-05
Epoch 81/100
814/814 [=====] - 3s 3ms/step - loss: 0.3230 -
val_loss: 0.3213 - lr: 1.0000e-05
Epoch 82/100
814/814 [=====] - 2s 3ms/step - loss: 0.3256 -
val_loss: 0.3223 - lr: 1.0000e-05
Epoch 83/100
814/814 [=====] - 2s 3ms/step - loss: 0.3233 -
val_loss: 0.3198 - lr: 1.0000e-05
Epoch 84/100
814/814 [=====] - 2s 3ms/step - loss: 0.3244 -
val_loss: 0.3201 - lr: 1.0000e-05
Epoch 85/100
814/814 [=====] - 3s 3ms/step - loss: 0.3245 -
val_loss: 0.3209 - lr: 1.0000e-05
Epoch 86/100
814/814 [=====] - 3s 4ms/step - loss: 0.3244 -
val_loss: 0.3198 - lr: 1.0000e-05
Epoch 87/100
814/814 [=====] - 3s 3ms/step - loss: 0.3227 -
val_loss: 0.3203 - lr: 1.0000e-05
Epoch 88/100
814/814 [=====] - 3s 3ms/step - loss: 0.3235 -
val_loss: 0.3217 - lr: 1.0000e-05
Epoch 89/100
814/814 [=====] - 3s 4ms/step - loss: 0.3250 -
val_loss: 0.3197 - lr: 1.0000e-05
Epoch 90/100
814/814 [=====] - 3s 3ms/step - loss: 0.3236 -
val_loss: 0.3193 - lr: 1.0000e-05
Epoch 91/100
814/814 [=====] - 3s 3ms/step - loss: 0.3244 -
val_loss: 0.3199 - lr: 1.0000e-05

```
[22]: train_loss = hist.history['loss']
      val_loss = hist.history['val_loss']
      #train_acc = hist.history['acc']
      #val_acc = hist.history['val_acc']

      import matplotlib.pyplot as plt

      plt.plot(train_loss,label='train loss')
      plt.plot(val_loss,label='val loss')
      plt.legend()
      plt.show()
```



```
[ ]:
```