

# Convolutional neural networks

## Practice Session 7

Changho Suh

January 24, 2024

# Recap: Two building blocks of CNNs

---

## 1. Convolutional layer (Conv layer)

**Role:** Mimick neurons' behaviors:  
Reacting only to receptive fields.

## 2. Pooling layer

**Role:** Downsample to reduce complexity  
(# parameters & memory size).

# Recap: Conv layer

---

1. A neuron reacts only to a **receptive field**.
2. A neighboring neuron concerns a receptive field shifted by **stride**.
3. Two types of **zero padding**: same, valid
4. Consists of a 2D **feature map**
5. Use of **multiple filters**
6. Has three **RGB** channels for colored images

# Recap: Pooling layer


---

1. Two types of pooling:  
Max pooling, average pooling
2. A default choice:  $2 \times 2$  pooling filter, stride=2
3. Works on each channel independently

# Recap: Two popular CNNs

## 1. **AlexNet** (2012)

Won the ImageNet competition.



Alex Krizhevsky Ilya Sutskever Geoffrey Hinton

Anchored the deep learning revolution.

## 2. **ResNet** (2015)

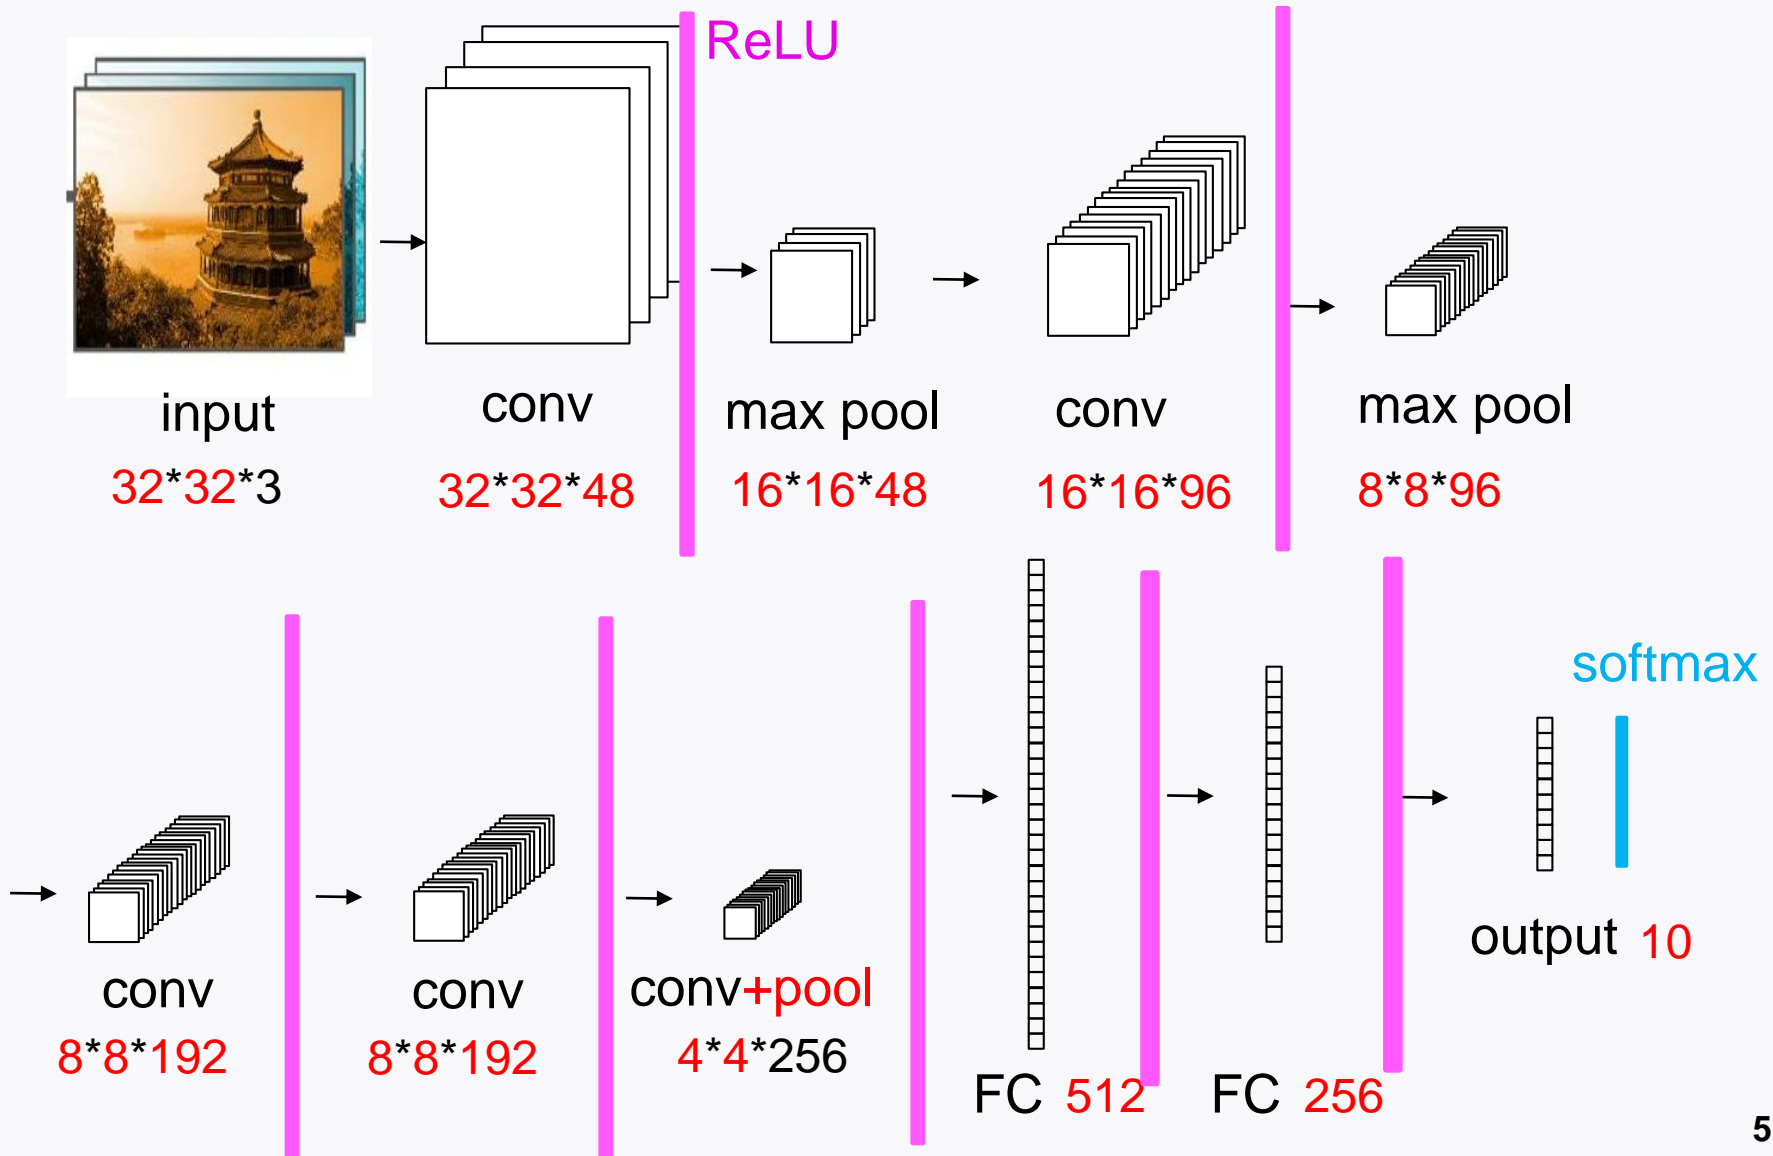
Won the 2015 ImageNet competition.

Currently the most powerful & arguably the simplest!

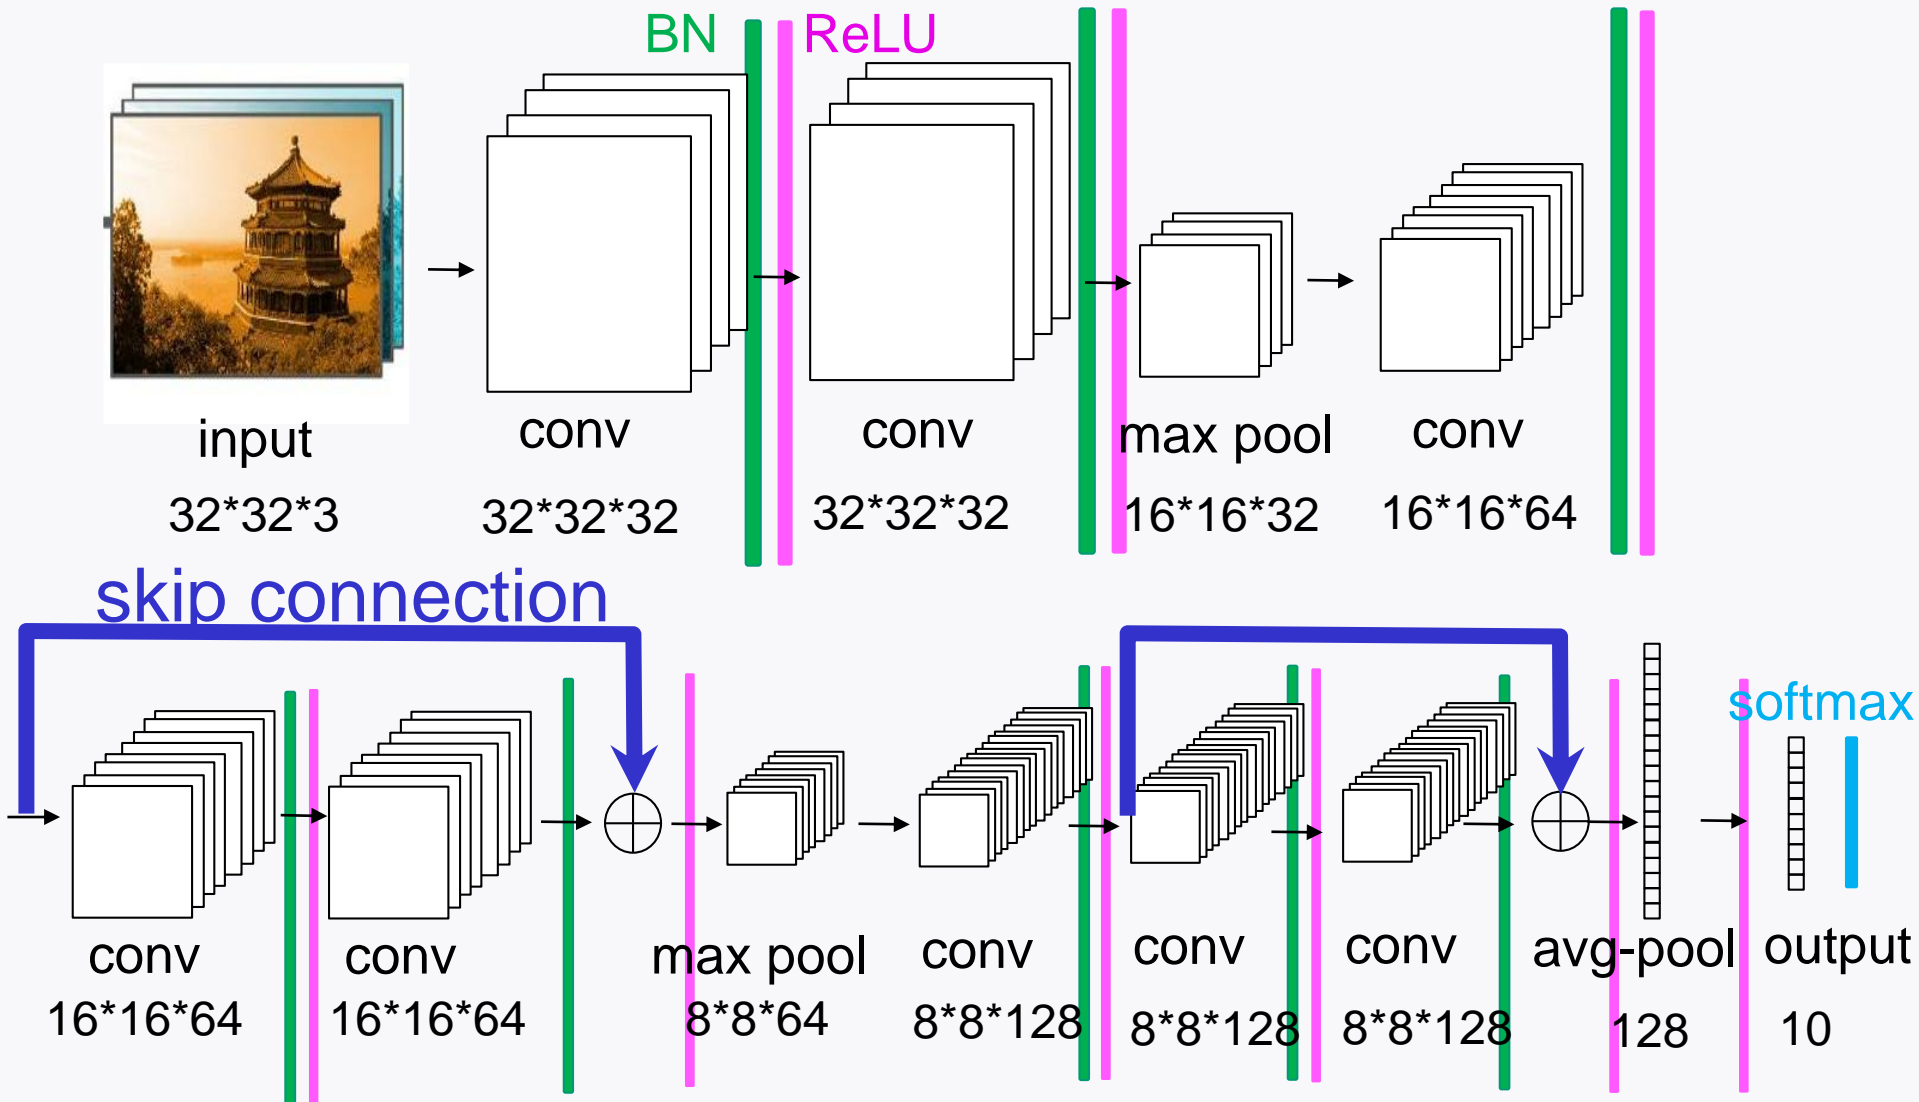


Kaiming He

# Recap: Simplified AlexNet



# Recap: Simplified ResNet



# Outline

---

Will implement three prominent models:

1. LeNet5

Task: Handwritten digit classification (MNIST)

2. Simplified AlexNet


Task: Image recognition (CIFAR10)

3. Simplified ResNet

Task: Image recognition (CIFAR10)




# Handwritten digit classification



# MNIST dataset


$$\{(x^{(i)}, y^{(i)})\}_{i=1}^m \quad m = 60,000$$

Examples:



Yann LeCun 1998


# Input image & label



# How to load MNIST dataset

```
from tensorflow.keras.datasets import mnist  
  
(X_train, y_train), (X_test, y_test) = mnist.load_data()  
  
X_train, X_test = X_train/255., X_test/255
```

# LeNet-5



# LeNet-5: Tensorflow coding

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Flatten, Dense


from tensorflow.keras.layers import Conv2D
from tensorflow.keras.layers import MaxPool2D
```

```
model_lenet = Sequential()
```

```
#1st stack ([Conv]+[ReLU]+[Pool])
```

```
model_lenet.add(Conv2D(input_shape=(28,28,1),
                        kernel_size=(5,5),
                        strides=(1,1),
                        filters=32,
                        padding='same',
                        activation='relu'
                    ))
```

```
model_lenet.add(MaxPool2D(pool_size=(2,2),
                           strides=(2,2),
                           padding='valid'))
```

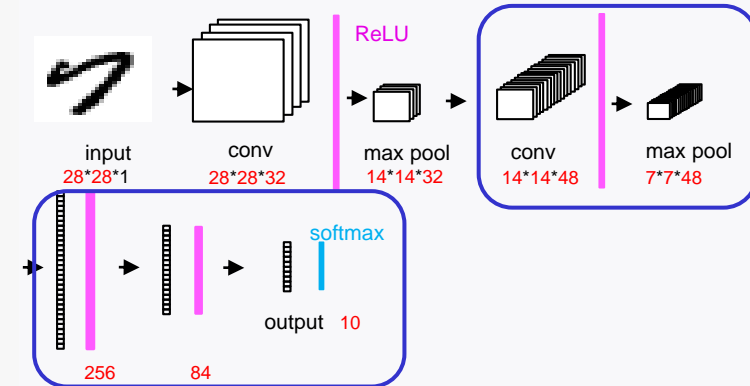


# LeNet-5: Tensorflow coding

*#2nd stack ([Conv]+[ReLU]+[Pool])*

```
model_lenet.add(Conv2D(kernel_size=(5,5),
                        strides=(1,1),
                        filters=48,
                        padding='same',
                        activation='relu'
                        ))

model_lenet.add(MaxPool2D(pool_size=(2,2),
                           strides=(2,2),
                           padding='valid'))
```



*# Three fully Connected Layers*

```
model_lenet.add(Flatten())
model_lenet.add(Dense(256,activation='relu' ))
model_lenet.add(Dense(84,activation='relu'))
model_lenet.add(Dense(10,activation='softmax'))
```

# Compile

```
from tensorflow.keras.optimizers import Adam

opt = Adam(learning_rate = 0.001,
           beta_1 = 0.9,
           beta_2 = 0.999)

model_lenet.compile(optimizer = opt,
                    loss='sparse_categorical_crossentropy',
                    metrics=['acc'])
```



# Training & evaluation

```
# Conversion from 3D to 4D tensor
X_train, X_test = X_train.reshape(-1,28,28,1), X_test.reshape(-1,28,28,1)

# Training
model_lenet.fit(X_train,y_train,epochs=10)

# Evaluation
test_performance = model_lenet.evaluate(X_test, y_test)
print(test_performance)
```

```
[0.036696918308734894, 0.9922999739646912]
```

# Look ahead

---

Will implement simplified AlexNet in the context of image recognition (CIFAR10).