

PS12

January 24, 2023

0.1 Recurrent Neural Networks

```
[1]: import tensorflow as tf
      from tensorflow.keras.datasets import imdb
      import numpy as np
```

0.2 Loading IMDB

```
[2]: from tensorflow.keras.datasets import imdb

      (X_train,y_train), (X_test,y_test) = imdb.load_data(num_words=10000)
```

```
[3]: # Get IMDB dictionary
      word_idx_dict = imdb.get_word_index()

      #print(word_idx_dict.keys())
      #print(word_idx_dict.values())
```

0.3 Data preprocessing: Padding

```
[4]: from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
[5]: print(len(X_train[0]))
```

218

```
[6]: max_len = 256
      X_train_pad = pad_sequences(X_train,value=0, padding='post', maxlen=256)
```

```
[7]: print(X_train_pad.shape)
      print(len(X_train_pad[0]))
      print(X_train_pad[0].shape)
      print(X_train_pad[0])
```

(25000, 256)

256

(256,)

[1 14 22 16 43 530 973 1622 1385 65 458 4468 66 3941

```

    4  173   36  256    5   25  100   43  838  112   50  670    2    9
   35  480  284    5  150    4  172  112  167    2  336  385   39    4
  172 4536 1111   17  546   38   13  447    4  192   50   16    6  147
2025   19   14   22    4 1920 4613  469    4   22   71   87   12   16
   43  530   38   76   15   13 1247    4   22   17  515   17   12   16
  626   18    2    5   62  386   12    8  316    8  106    5    4 2223
5244   16  480   66 3785   33    4  130   12   16   38  619    5   25
  124   51   36  135   48   25 1415   33    6   22   12  215   28   77
   52    5   14  407   16   82    2    8    4  107  117 5952   15  256
    4    2    7 3766    5  723   36   71  43  530  476   26  400  317
   46    7    4    2 1029   13  104   88    4  381   15  297   98   32
2071   56   26  141    6  194 7486   18    4  226   22   21  134  476
   26  480    5  144   30 5535   18   51   36   28  224   92   25  104
    4  226   65   16   38 1334   88   12   16  283    5   16 4472  113
  103   32   15   16 5345   19  178   32    0    0    0    0    0    0
    0    0    0    0    0    0    0    0    0    0    0    0    0    0
    0    0    0    0    0    0    0    0    0    0    0    0    0    0
    0    0    0    0]

```

```
[8]: #word_idx_dict_new['<PAD>']=0
      #word_idx_dict_new
```

```
[9]: #print(decode_review(X_train_pad[0]))
```

```
[10]: X_test_pad = pad_sequences(X_test, value=0, padding='post',maxlen=max_len)
```

```
[11]: print(X_train_pad.shape)
      print(X_test_pad.shape)
```

```
(25000, 256)
```

```
(25000, 256)
```

0.4 Simple RNN

```
[13]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Dense
      from tensorflow.keras.layers import Embedding
      from tensorflow.keras.layers import SimpleRNN

      basic_rnn = Sequential()
      basic_rnn.add(Embedding(input_dim=10000, output_dim=100, input_length=256))
      basic_rnn.add(SimpleRNN(128))
      basic_rnn.add(Dense(1,activation='sigmoid'))
```

```
[13]: # validation dataset
      X_train_pad, X_val_pad = X_train_pad[:20000], X_train_pad[20000:]
      y_train, y_val = y_train[:20000], y_train[20000:]
```

```
[14]: basic_rnn.compile(optimizer='adam',loss='binary_crossentropy',metrics='acc')
      #history_basic_rnn = basic_rnn.fit(X_train_pad,y_train,
      #                                validation_data=(X_val_pad,y_val),
      #                                epochs=10)
      history_basic_rnn = basic_rnn.fit(X_train_pad,y_train, epochs=10)
```

```
Epoch 1/10
782/782 [=====] - 24s 30ms/step - loss: 0.6980 - acc:
0.5087
Epoch 2/10
782/782 [=====] - 25s 32ms/step - loss: 0.6946 - acc:
0.5164
Epoch 3/10
782/782 [=====] - 25s 33ms/step - loss: 0.6777 - acc:
0.5491
Epoch 4/10
782/782 [=====] - 26s 33ms/step - loss: 0.6458 - acc:
0.5723
Epoch 5/10
782/782 [=====] - 27s 35ms/step - loss: 0.6249 - acc:
0.6053
Epoch 6/10
782/782 [=====] - 27s 34ms/step - loss: 0.6476 - acc:
0.5698
Epoch 7/10
782/782 [=====] - 27s 34ms/step - loss: 0.6417 - acc:
0.5758
Epoch 8/10
782/782 [=====] - 26s 34ms/step - loss: 0.6218 - acc:
0.5865
Epoch 9/10
782/782 [=====] - 27s 34ms/step - loss: 0.5929 - acc:
0.6069
Epoch 10/10
782/782 [=====] - 27s 34ms/step - loss: 0.5744 - acc:
0.6210
```

```
[15]: test_performance = basic_rnn.evaluate(X_test_pad,y_test)
      print(test_performance)
```

```
782/782 [=====] - 8s 10ms/step - loss: 0.7149 - acc:
0.5500
[0.7149398922920227, 0.5500400066375732]
```

0.5 LSTM

```
[23]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Dense
      from tensorflow.keras.layers import Embedding
      from tensorflow.keras.layers import LSTM

      LSTM_model = Sequential()
      LSTM_model.add(Embedding(input_dim=10000,output_dim=100,input_length=256))
      LSTM_model.add(LSTM(128))
      LSTM_model.add(Dense(1,activation='sigmoid'))
      LSTM_model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 256, 100)	1000000
lstm (LSTM)	(None, 128)	117248
dense_2 (Dense)	(None, 1)	129

=====
Total params: 1,117,377
Trainable params: 1,117,377
Non-trainable params: 0
=====

```
[24]: LSTM_model.compile(optimizer='adam',
                        loss='binary_crossentropy',
                        metrics='acc')

      LSTM_model.fit(X_train_pad,y_train, epochs=10)
```

Epoch 1/10
625/625 [=====] - 91s 144ms/step - loss: 0.6758 - acc: 0.5514
Epoch 2/10
625/625 [=====] - 87s 140ms/step - loss: 0.6407 - acc: 0.5856
Epoch 3/10
625/625 [=====] - 88s 141ms/step - loss: 0.6165 - acc: 0.6079
Epoch 4/10
625/625 [=====] - 86s 138ms/step - loss: 0.5761 - acc: 0.6523
Epoch 5/10

```

625/625 [=====] - 87s 139ms/step - loss: 0.4452 - acc:
0.7847
Epoch 6/10
625/625 [=====] - 85s 136ms/step - loss: 0.2666 - acc:
0.8983
Epoch 7/10
625/625 [=====] - 85s 136ms/step - loss: 0.1760 - acc:
0.9370
Epoch 8/10
625/625 [=====] - 86s 138ms/step - loss: 0.1161 - acc:
0.9640
Epoch 9/10
625/625 [=====] - 85s 137ms/step - loss: 0.0744 - acc:
0.9805
Epoch 10/10
625/625 [=====] - 86s 138ms/step - loss: 0.0536 - acc:
0.9872

```

[24]: <keras.callbacks.History at 0x17731d55280>

```

[76]: #test_performance = LSTM_model.evaluate(X_test_pad,y_test)[1]
test_performance = LSTM_model.evaluate(X_test_pad,y_test)
print(test_performance)

```

```

782/782 [=====] - 22s 28ms/step - loss: 0.5110 - acc:
0.8502
782/782 [=====] - 22s 28ms/step - loss: 0.5110 - acc:
0.8502
[0.5110498666763306, 0.8502399921417236]

```