

Small data technique

Practice Session 14

Changho Suh

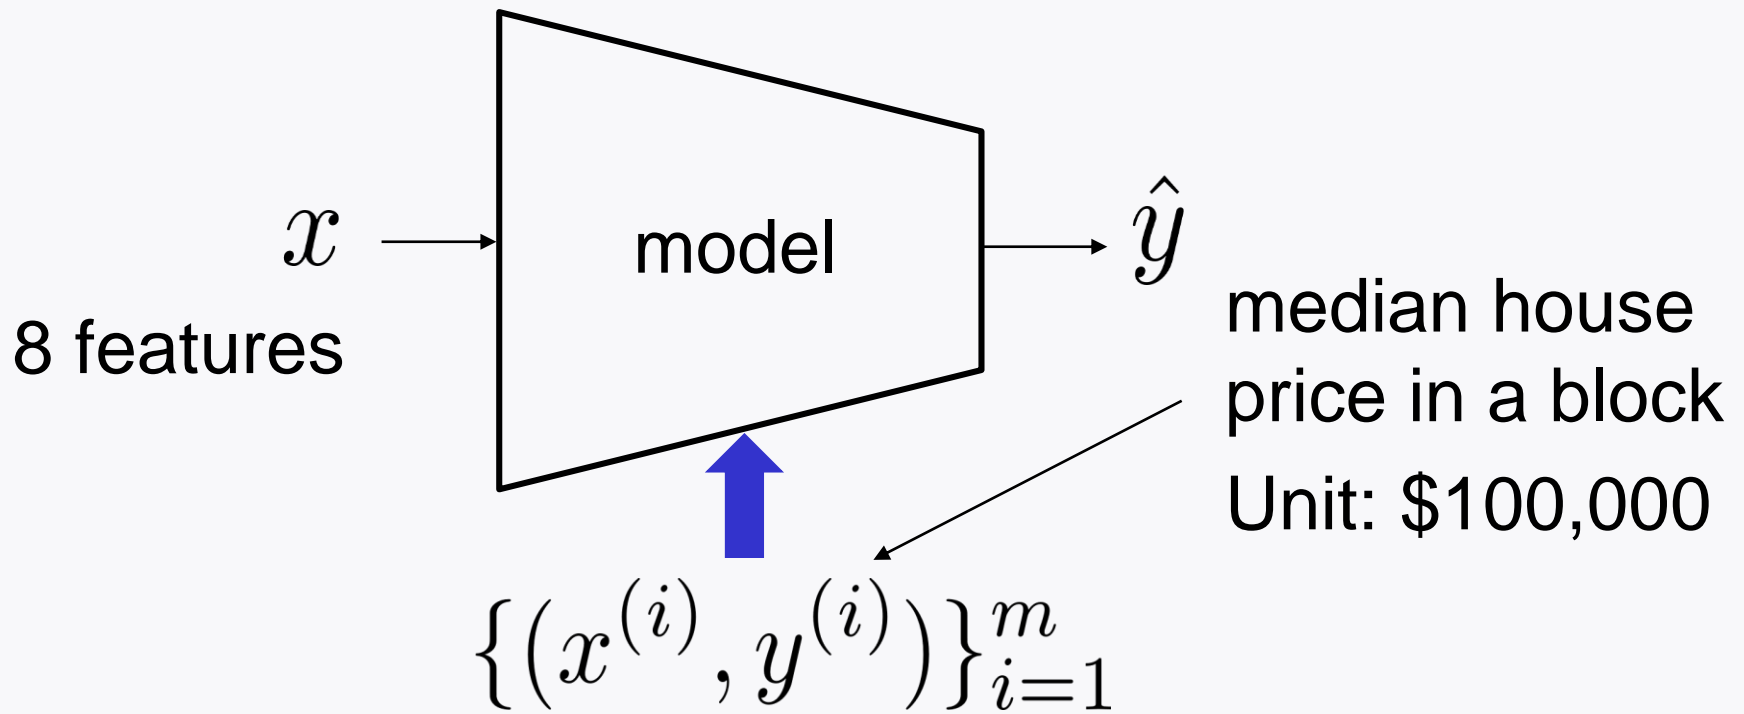
January 26, 2024

Outline

Implementation of **decision tree regressor**

Task: California housing price prediction

California housing price prediction



$$m = 20,640$$

8 features

MedInc	median income in block group
HouseAge	median house age in block group
AveRooms	average # of rooms per household
AveBedrms	average # of bedrooms per household
Population	block group population
AveOccup	average # of household members
Latitude	block group latitude
Longitude	block group longitude

Load California Housing dataset

```
from sklearn.datasets import fetch_california_housing
```

```
cali_prices = fetch_california_housing()
```

```
X_reg = cali_prices.data
```

```
y_reg = cali_prices.target
```

```
print(X_reg.shape) (20640, 8)
```

```
print(y_reg.shape) (20640,)
```

```
print(cali_prices.feature_names)
```

```
['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup', 'Latitude',  
'Longitude']
```

Data

```
['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup', 'Latitude',  
'Longitude']
```

```
print(X_reg[0])
```

```
[ 8.3252      41.      6.98412698  1.02380952 322.  
 2.55555556 37.88    -122.23      ]
```

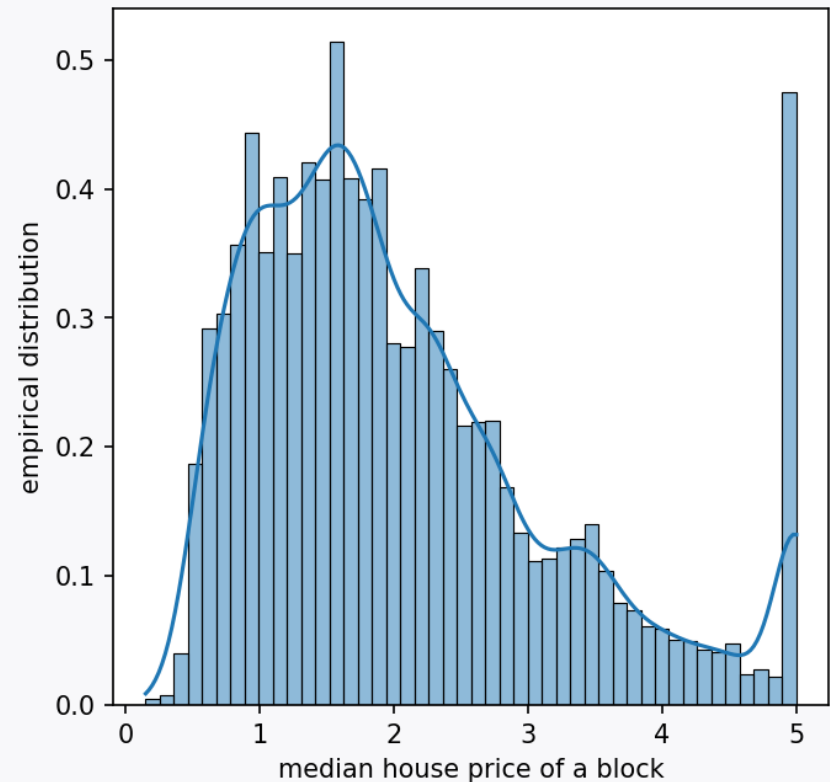
```
print(y_reg)
```

```
[4.526 3.585 3.521 ... 0.923 0.847 0.894]
```

Statistics of median housing price

```
import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(5,5),dpi=150)
sns.histplot(y_reg, kde=True, stat='density')
plt.xlabel('median house price of a block')
plt.ylabel('empirical distribution')
plt.show()
```



Decision tree regressor

```
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor

# train-test data split
X_train, X_test, y_train, y_test = train_test_split(X_reg, y_reg, test_size=0.01)

tree_reg = DecisionTreeRegressor(max_depth=4)

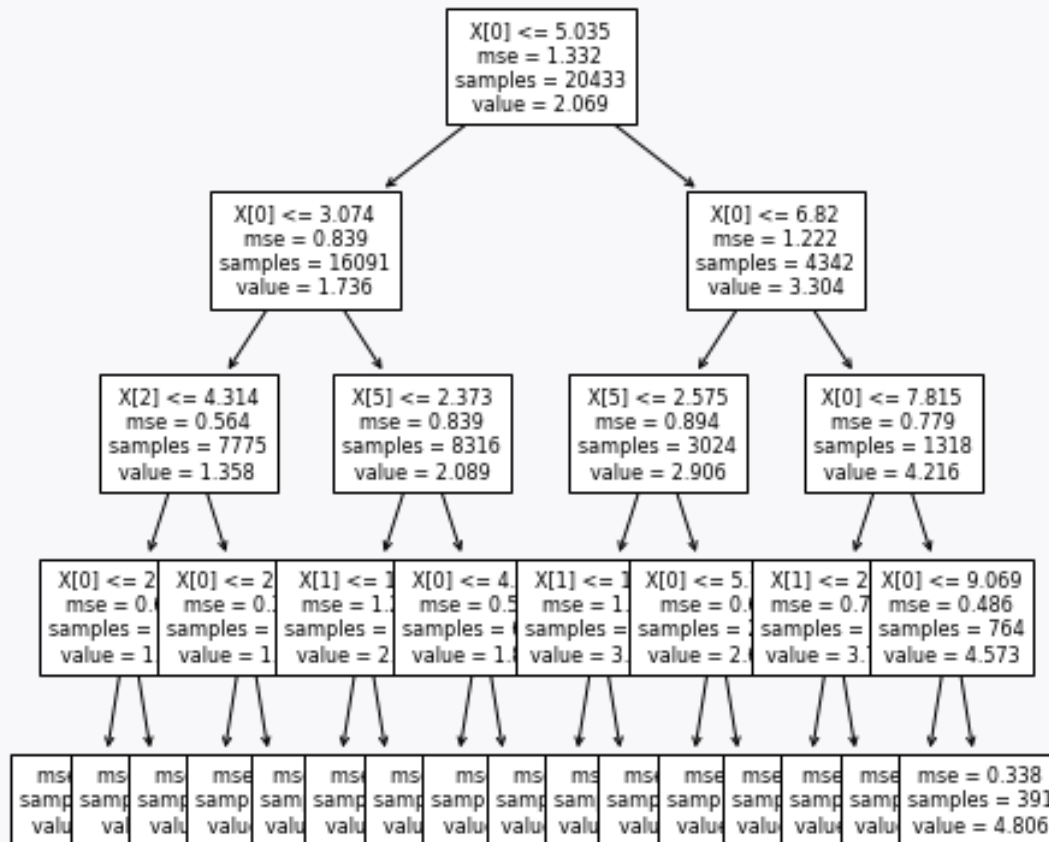
# training
tree_reg.fit(X_train, y_train)

# evaluation (R2 score)
test_performance = tree_reg.score(X_test, y_test)
print(test_performance)
```

```
0.6985944550597363
```


Regressor visualization

```
from sklearn.tree import plot_tree
plot_tree(tree_reg, fontsize=8)
```



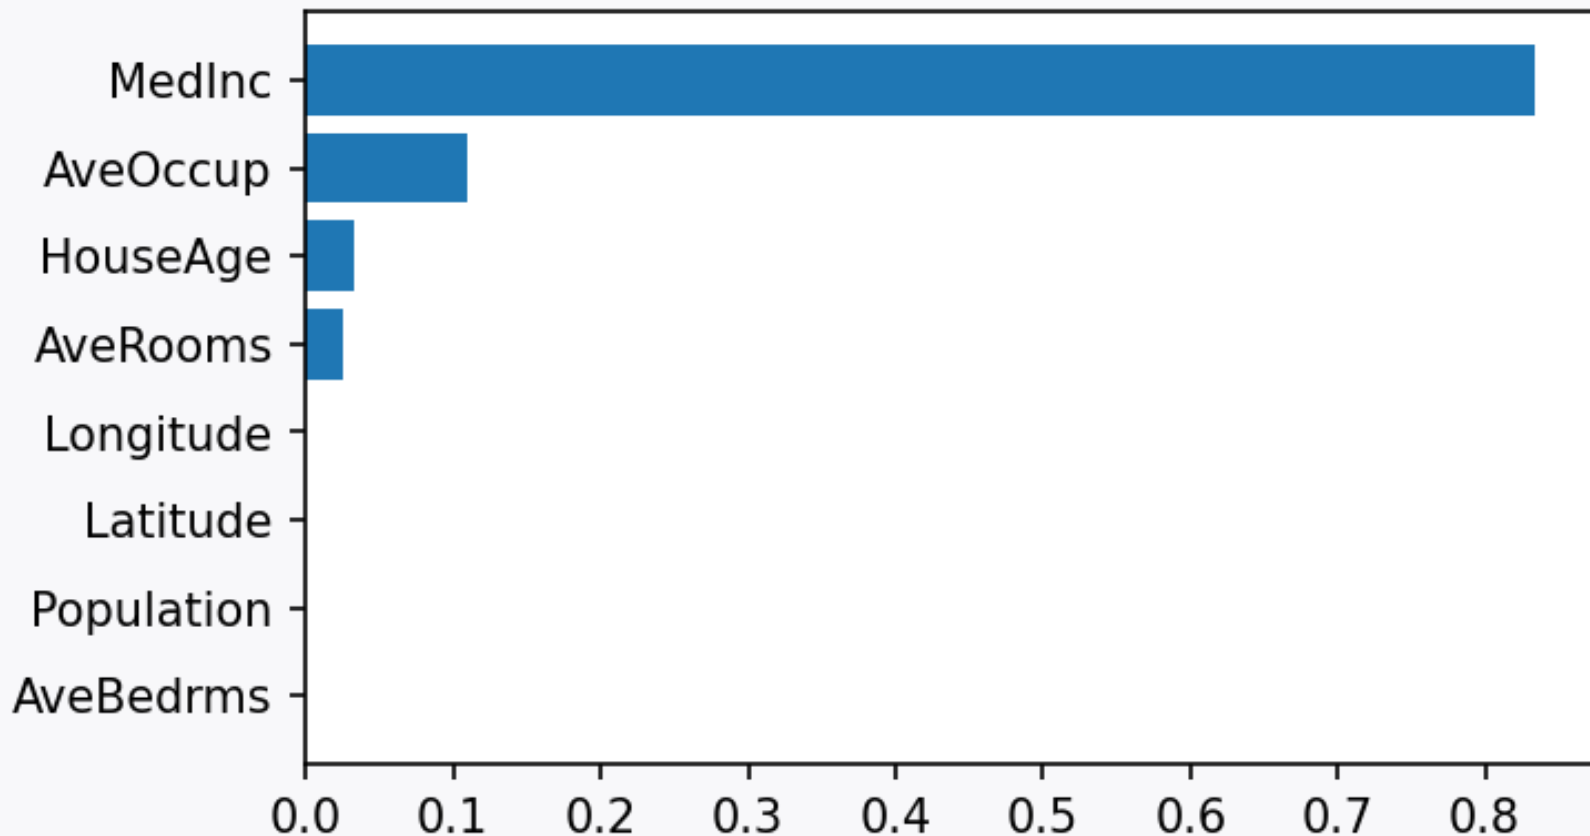
Feature importance

```
feature_importances_cali = tree_reg.feature_importances_  
print(feature_importances_cali)
```

```
[0.83298645 0.03137747 0.02509265 0.          0.          0.11054343  
 0.          0.          ]
```

Feature importance visualization

```
sorted_idx = tree_reg.feature_importances_.argsort()  
plt.barh(np.asarray(cali_prices.feature_names)[sorted_idx],  
         feature_importances_cali[sorted_idx])
```



MSE performance

```
from sklearn.metrics import mean_squared_error
```

```
y_pred_train = tree_reg.predict(X_train)  
y_pred_test = tree_reg.predict(X_test)
```

```
mse_train = mean_squared_error(y_pred_train, y_train)  
mse_test = mean_squared_error(y_pred_test, y_test)  
print(mse_train)  
print(mse_test)
```

```
0.5556284248376379  
0.3857388225897835
```

Comparison with random guess

```
random_guess_train = np.random.normal(  
    loc = np.mean(y_train),  
    scale = np.std(y_train),  
    size = y_train.shape[0])
```

```
random_guess_test = np.random.normal(  
    loc = np.mean(y_test),  
    scale = np.std(y_test),  
    size = y_test.shape[0])
```

```
mse_first_train = mean_squared_error(random_guess_train, y_train)  
mse_first_test = mean_squared_error(random_guess_test, y_test)  
print(mse_first_train)  
print(mse_first_test)
```

```
2.6862254154194662  
2.6724862045770768
```

Look ahead

Implementation of **random forests**

Tasks: Iris plants classification

California housing price prediction

Handwritten digit classification